Meta-Learning for Generalized Zero-Shot Learning

Vinay Kumar Verma, Dhanajit Brahma, Piyush Rai
Department of Computer Science and Engineering, IIT Kanpur, India
{vkverma,dhanajit,piyush}@cse.iitk.ac.in

Abstract

Learning to classify unseen class samples at test time is popularly referred to as zero-shot learning (ZSL). If test samples can be from training (seen) as well as unseen classes, it is a more challenging problem due to the existence of strong bias towards seen classes. This problem is generally known as *generalized* zero-shot learning (GZSL). Thanks to the recent advances in generative models such as VAEs and GANs, sample synthesis based approaches have gained considerable attention for solving this problem. These approaches are able to handle the problem of class bias by synthesizing unseen class samples. However, these ZSL/GZSL models suffer due to the following key limitations: (i) Their training stage learns a class-conditioned generator using only *seen* class data and the training stage does not *explicitly* learn to generate the unseen class samples; (ii) They do not learn a generic optimal parameter which can easily generalize for both seen and unseen class generation; and (iii) If we only have access to a very few samples per seen class, these models tend to perform poorly. In this paper, we propose a meta-learning based generative model that naturally handles these limitations. The proposed model is based on integrating model-agnostic meta learning with a Wasserstein GAN (WGAN) to handle (i) and (iii), and uses a novel task distribution to handle (ii). Our proposed model yields significant improvements on standard ZSL as well as more challenging GZSL setting. In ZSL setting, our model yields 4.5%, 6.0%, 9.8%, and 27.9% relative improvements over the current state-of-the-art on CUB, AWA1, AWA2, and aPY datasets, respectively.

Introduction

With the ever-growing quantities, diversity, and complexity of real-world data, machine learning algorithms are increasingly faced with challenges that are not adequately addressed by traditional learning paradigms. For classification problems, one such challenging setting is where test-time requires correctly labeling objects that could be from classes that were not present at training time. This setting is popularly known as Zero-Shot Learning (ZSL), and has drawn a considerable interest recently (Socher et al. 2013; Norouzi et al. 2013; Verma and Rai 2017; Changpinyo et al. 2016;

Romera and Torr 2015a; Xian et al. 2018b; Verma et al. 2018; Liu et al. 2018; Romera and Torr 2015b; Chen et al. 2018). ZSL algorithms typically rely on class-descriptions (e.g., human-provided class attribute vectors, textual description, or word2vec embedding of class name). These class-description/class-attributes are leveraged to transfer the knowledge from *seen* classes (i.e., classes that were present at training-time) to *unseen* classes (i.e., classes only encountered in test data).

Driven by the recent advances in generative modeling (Arjovsky, Chintala, and Bottou 2017; Kingma and Welling 2014), there is a growing interest in generative models for ZSL. Broadly, these models learn to generate/synthesize "artificial" examples from unseen classes (Felix et al. 2018; Verma et al. 2018; Xian et al. 2018b), conditioning on their class attributes, and learn a classifier using these synthesized examples. Despite the recent progress on such approaches, these still have some key limitations. Firstly, while the goal of these approaches is to generate unseen/novel class examples given the respective class attributes, these models are trained using data (inputs and the respective class attributes) from the seen classes (Verma et al. 2018; Xian et al. 2018b; Felix et al. 2018) and do not explicitly learn to generate the unseen class samples during training. Consequently, these generative ZSL models show a large quality gap between the synthesized unseen class inputs and actual unseen class input. To mimic the ZSL setting explicitly, we propose a novel variant of the standard meta-learning based approach (Finn, Abbeel, and Levine 2017). Notably, in our variant, the meta-train and meta-validation classes are *disjoint*.

The second limitation of existing ZSL/GZSL models is that they do not learn an optimal parameter which can easily generalize to the seen/unseen class generation. Our meta-learning framework learns such an optimal parameter that can quickly adapt to the novel classes (meta-test) with few gradient steps. (Snell, Swersky, and Zemel 2017; Vinyals et al. 2016) show that even with the *zero-gradient step* (without fine-tuning), meta-learning learns to generalize novel class samples/task. We build on this idea to train a class-conditioned WGAN for sample generation.

The third key limitation is that all the existing ZSL methods rely on the availability of a significant number of labeled

samples from each of the seen classes. This itself is a severe requirement and may not be met in practice (e.g., we may only have a handful, say 5, or 10 examples from each seen class). Note that this setting is somewhat similar to few-shot learning or meta-learning (Finn, Abbeel, and Levine 2017) where the goal is to learn a classifier using very few examples per class, but all the test/unseen class are assumed to have few samples in test time. In contrast, in ZSL, we do not have any labeled training data from unseen classes. Our meta-learning based formulation is naturally suited to this setting where only a few samples per class are available.

Our approach is primarily based on learning a generative model that can synthesize inputs from any class (seen/unseen), given the respective class-attributes/description. However, unlike recent works on synthesis based ZSL models (Zhu et al. 2018; Verma et al. 2018; Xian et al. 2018b; Felix et al. 2018), we endow the generator the capability to meta-learn using very few examples per seen class. To this end, we develop a meta-learning based *conditional* Wasserstein GAN (Arjovsky, Chintala, and Bottou 2017) (conditioning on the class-attributes) which has a generator and a discriminator modules augmented with a classifier. Each module is associated with a meta-learning agent, to facilitate learning with a very small number of seen class inputs. Also, the novel task distribution helps to mimic the ZSL behavior, i.e., the generative model not only learns to generate the seen class samples but the unseen class samples as well. We would also like to highlight that, although we develop this model with the focus being ZSL and generalized ZSL, our ideas can be used for the task of supervised few-shot generation (Clouâtre and Demers 2019), which is the problem of learning to generate data given very few examples to learn the data distribution. Our main contributions are summarized below:

- We develop a novel meta-learning framework for ZSL and generalized ZSL by learning to synthesize examples from unseen classes, given the respective class-attributes. Notably, our framework is based on model-agnostic meta-learning (Finn, Abbeel, and Levine 2017), which enables the synthesis of high-quality examples. This helps to overcome the above mentioned second and third limitation.

- We propose a novel episodic training for the meta-learning based ZSL where, in each episode, the *training-set* and *validation-set* classes are disjoint. This helps *learning to generate the novel class examples in training itself*. This contributes in overcoming the above mentioned first limitation.

Notation, Preliminaries, Problem Setup

A typical ZSL setting is as follows: We have S *seen* classes with labelled training data and U *unseen* classes with no labelled data present during the training time. The test data can be either exclusively from unseen classes (standard ZSL setting), or can be from both unseen and seen classes (*generalized* ZSL setting). We further assume that we are provided *class-attribute vectors* for the seen as well as unseen classes $\mathbf{A} = \{\mathbf{a_c}\}_{c=1}^{S+U}$, where $\mathbf{a_c} \in \mathbb{R}^d$ is the class-attribute vector of class c. These class-attribute vectors are leveraged by the ZSL algorithms to transfer the knowledge from seen to unseen classes.

Existing ZSL algorithms 1assume that we have access to a significant number of examples from each of the seen classes. This may however not be the case; in practice, we may have very few examples from each of the seen classes. We train our model in N-way K-shot setting such that it can handle the ZSL problem when only very few samples are available per seen class. We choose the model-agnostic meta learning (MAML) (Finn, Abbeel, and Levine 2017) as our meta-learner due to its generic nature; it only requires a differentiable model and can work with any loss function.

Model-Agnostic Meta-Learning (MAML)

MAML (Finn, Abbeel, and Levine 2017) is an optimization based meta-learning framework designed for few-shot learning. The model is designed in such a way that it can quickly adapt to a new task with the help of only few training examples. MAML assumes that model f_θ is parameterized by learnable parameters θ and the loss function is smooth in θ that can be used for the gradient-descent based updates.

Let $p(\mathcal{T})$ be the distribution of tasks over the *meta-train* set. MAML defines the notion of a "task" such that a task $\mathcal{T}_i \sim p(\mathcal{T})$ represents a set of labeled examples and MAML splits this set further into a training set $\mathcal{T}_{tr}$ and a validation set $\mathcal{T}_{val}$, i.e., $\mathcal{T}_i = \{\mathcal{T}_{tr}, \mathcal{T}_{val}\}$. The split is done such that $\mathcal{T}_{tr}$ has very few examples per class. We follow the general notion of N-way K-shot problem (Vinyals et al. 2016), i.e., $\mathcal{T}_{tr}$ contains N classes with K examples from each class. The model is trained using an episodic formulation where each round samples a batch of tasks and uses gradient-descent based updates (inner loop) for the parameters θ_i specific to each task $\mathcal{T}_i$. The meta-update step (outer loop) then aggregates the information from all these "local" updates to update the overall model parameters θ, using gradient descent update.

For task $\mathcal{T}_i$, its local parameters θ_i are updated by starting with the global model parameters θ, and using a few gradient based updates computed on $\mathcal{T}_{tr}$ from task $\mathcal{T}_i$. Assuming a single step of update, this can be written as: $\theta_i' = \theta - \alpha \nabla_\theta \mathcal{L}_{\mathcal{T}_{tr}}(f_\theta)$. Here, α is the hyper-parameter and $\mathcal{L}$ denotes the loss function being used. The overall global/meta objective defined over the multiple tasks sampled from task distribution $p(\mathcal{T})$ can be defined as:

$$\sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_{tr}}(f_{\theta_i'}) = \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_{tr}}(f_{\theta - \alpha \mathcal{L}_{\mathcal{T}_{tr}}(f_\theta)}) \qquad (1)$$

Assuming a gradient descent based optimization of the global objective in Eq. 1, a single-step gradient descent update for the global parameter can be written as: $\theta \leftarrow \theta - \beta \nabla_\theta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_{val}}(f_{\theta_i'})$.

Zero-Shot Meta-Learning (ZSML)

The meta-learning framework can quickly adapt to a new task with the help of only a few gradient steps. The quick adaption is only possible for the model if it learns the optimal parameter θ in the parameter space that is unbiased towards the meta-train data. The learned parameters are close
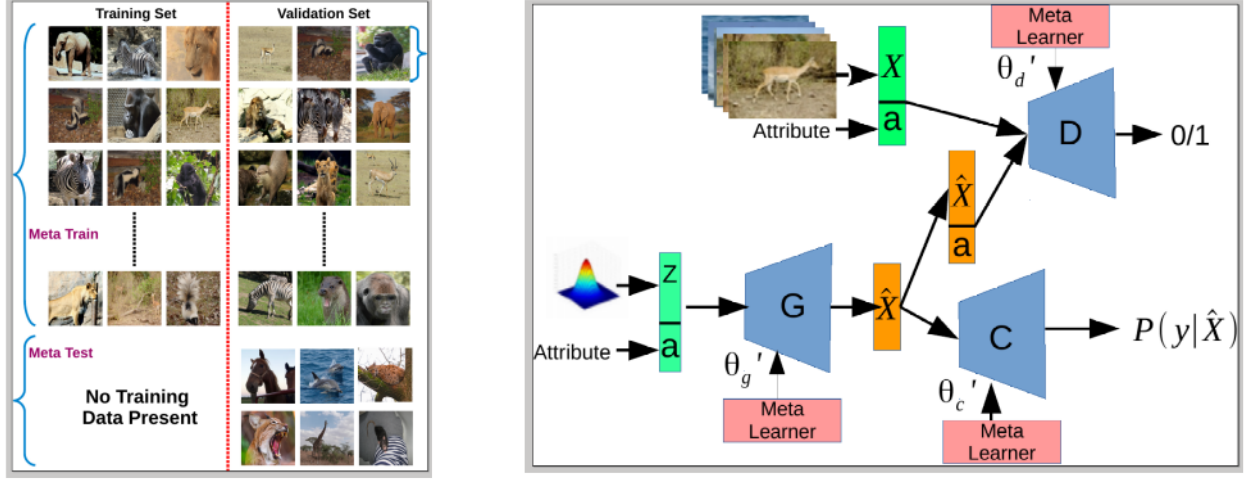
Figure 1: **Left:** Task episode for zero-shot meta-learning. For each task $\mathcal{T}_i = \{\mathcal{T}_{tr}, \mathcal{T}_{val}\}$, *training set $\mathcal{T}_{tr}$* and *validation set $\mathcal{T}_{val}$* classes are disjoint. In the ZSL setup, we have *zero* training examples from the *meta-test* set. **Right:** The proposed architecture model. X: ResNet-101 feature vector.

to the optimal parameters for both meta-train and meta-test data respectively. It is already demonstrated in (Vinyals et al. 2016; Snell, Swersky, and Zemel 2017) where without fine-tuning (using *zero* gradient steps, i.e., not making any update) on the meta-test, the meta-learning model shows better/similar performance. Our ZSML approach is primarily motivated by high-quality generalization ability of the meta-learning towards the seen/unseen class samples. We use the meta-learning framework to train a generative adversarial network conditioned on class attributes, that can generate the novel class samples. A key difference with MAML, to *mimic the ZSL behaviour*, is that for each task $\mathcal{T}_i = \{\mathcal{T}_{tr}, \mathcal{T}_{val}\}$, the classes of $\mathcal{T}_{tr}$ and $\mathcal{T}_{val}$ are disjoint, whereas, in MAML, both set of classes are the same. Therefore, the training is done in such a way that $\mathcal{T}_{tr}$ acts as *seen* classes and $\mathcal{T}_{val}$ acts as *unseen* classes. The inner loop of the meta-learning optimizes the parameters using $\mathcal{T}_{tr}$, and final parameters are updated over the loss of the $\mathcal{T}_{val}$ (containing disjoint set of classes). Therefore, the model *learns to generate the novel class during the training itself*. In the next section, we describe our complete model (shown in Figure 1 (right)).

Meta-Learning based Adversarial Generation

The core of our ZSL model (Figure 1 right) is a generative adversarial network (Goodfellow et al. 2014), coupled with (1) an additional classifier module trained to correctly classify the examples generated by the generator module; and (2) meta-learners in each of the three modules (Generator (G), Discriminator (D), and Classifier (C)). We use the Wasserstein GAN (Arjovsky, Chintala, and Bottou 2017) architecture due to its nice stability properties. We assume θ_d, θ_g and θ_c to be the parameters of the Discriminator, Generator and Classifier, respectively.

Our model follows the episode-wise training akin to MAML (however, $\mathcal{T}_{tr}$ and $\mathcal{T}_{val}$ classes are disjoint in our ZSL setting). There are three meta-learners in the model, one for each D, G and C, but G and C are optimized jointly. From now on, we will denote the parameters for G and C as a joint set of parameters $\theta_{gc} = [\theta_g, \theta_c]$.

For each task $\mathcal{T}_i = \{\mathcal{T}_{tr}, \mathcal{T}_{val}\}$, sampled from the task distribution $p(\mathcal{T})$, $\mathcal{T}_{tr}$ is used by the meta-learners (in the inner loop) of D and G. $\mathcal{T}_{val}$ is used to calculate the loss over the most recent parameters of the meta-learners. For our model, the generator network $G : \mathbf{Z} \times \mathbf{A} \rightarrow \hat{\mathbf{X}}$ takes input as, a random noise $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ ($\mathbf{z} \in \mathbf{Z}$), concatenated with the class-attribute vector $\mathbf{a_c}$ of a class. G produces a sample $\hat{\mathbf{x}} \in \hat{\mathbf{X}}$ that is similar to a real sample from *that* class. The discriminator network $D : \mathbf{X} \times \mathbf{A} \rightarrow [0, 1]$ tries to distinguish such generated samples (concatenated with attributes) from the actual sample $\mathbf{X}$ (real data distribution). In addition, the goal of the classifier network $C : \hat{\mathbf{X}} \rightarrow \mathcal{Y}$ is to take the generated sample $\hat{\mathbf{x}}$ from G and classify it into the original class $c \in \mathcal{Y}$ where $\mathcal{Y}$ is the set of both *seen* and *unseen* classes. Presence of the classifier module C ensures that the generated sample has the same characteristics as that of samples from that class.

We now describe the objective function of our model. Let $\mathcal{L}_{\mathcal{T}_i}^D$ denote the meta-learner objective of the discriminator D and $\mathcal{L}_{\mathcal{T}_i}^{GC}$ denote the meta-learner objective of the generator G and the classifier C, on the task $\mathcal{T}_i$. The meta-learner objective $\mathcal{L}_{\mathcal{T}_i}^D$ for discriminator D can be defined as:

$$\mathcal{L}_{\mathcal{T}_i}^D(\theta_d) = \mathbb{E}_{\mathcal{T}_i} D(\mathbf{x}, \mathbf{a_c}|\theta_d) - \mathbb{E}_{\mathbf{a_c}, \hat{\mathbf{x}} \sim P_{\theta_g}} D(\hat{\mathbf{x}}, \mathbf{a_c}|\theta_d) \quad (2)$$

Here, $\mathbf{a_c} \in \mathbf{A}$ is attribute vector of samples belonging to $\mathcal{T}_i$. The objective in Eq. 2 (to be maximized) essentially says that the discriminator should have $D(.)$ large for real examples and small for generated examples. The meta-learner objective $\mathcal{L}_{\mathcal{T}_i}^{GC}$ for generator G and classifier C is given as:

$$\mathcal{L}_{\mathcal{T}_i}^{GC}(\theta_{gc}) = - \mathbb{E}_{\mathbf{a_c}, \mathbf{z} \sim \mathcal{N}(0, I)} D(G(\mathbf{a_c}, \mathbf{z}|\theta_g), \mathbf{a_c}|\theta_d) + C(y|\hat{\mathbf{x}}, \theta_c)$$
$$(3)$$

This objective (to be minimized) says that the generator's output $G(\mathbf{a_c}, \mathbf{z}|\theta_g)$ should be such that $D(.)$ is large, as well

as the classifier's loss C should be small (i.e., the classifier should predict the correct class for generated example $\hat{\mathbf{x}}$). Having defined the individual objectives, the overall objective for the meta-learner (inner loop) update for task $\mathcal{T}_i$:

$$l^D_{\mathcal{T}_i} = \max_{\theta_d} \mathcal{L}^D_{\mathcal{T}_i}(\theta_d) \quad \text{and} \quad l^{GC}_{\mathcal{T}_i} = \min_{\theta_{gc}} \mathcal{L}^{GC}_{\mathcal{T}_i}(\theta_{gc}) \quad (4)$$

The meta-learner gradient ascent update for the discriminator over a task $\mathcal{T}_i$ will be:

$$\theta'_d = \theta_d + \eta_1 \nabla_{\theta_d} l^D_{\mathcal{T}_{tr} \in \mathcal{T}_i}(\theta_d) \quad (5)$$

Similarly the meta-learner gradient descent update for the generator and classifier over $\mathcal{T}_i$ will be:

$$\theta'_{gc} = \theta_{gc} - \eta_2 \nabla_{\theta_{gc}} l^{GC}_{\mathcal{T}_{tr} \in \mathcal{T}_i}(\theta_{gc}) \quad (6)$$

The model parameters are learned by optimizing Eq 2 and Eq 3 over a batch of sampled tasks from the task distribution $p(\mathcal{T})$. The overall meta-objective for the discriminator and generator is:

$$\theta'_d = \theta_d + \eta_1 \nabla_{\theta_d} \sum_{\mathcal{T}_{tr} \in \mathcal{T}_i \sim p(\mathcal{T})} l^D_{\mathcal{T}_{tr}}(\theta_d) \quad (7)$$

$$\theta'_{gc} = \theta_{gc} - \eta_2 \nabla_{\theta_{gc}} \sum_{\mathcal{T}_{tr} \in \mathcal{T}_i \sim p(\mathcal{T})} l^{GC}_{\mathcal{T}_{tr}}(\theta_{gc}) \quad (8)$$

Unlike to standard MAML in the inner loop (i.e. Eq:7 and 8) are optimize on the set of task instead of per task. We observe that this increase the stability of the WGAN training. Having meta-learned the discriminator parameters from the meta-training phase (performed using the seen class examples), the discriminator's objective function w.r.t. the unseen class examples in the validation meta-set is given by:

$$\max_{\theta_d} \sum_{\mathcal{T}_{val} \in \mathcal{T}_i \sim p(\mathcal{T})} l^D_{\mathcal{T}_{val}}(\theta'_d)$$
$$= \max_{\theta_d} \sum_{\mathcal{T}_{val} \in \mathcal{T}_i \sim p(\mathcal{T})} l^D_{\mathcal{T}_{val}}(\theta_d + \eta_1 \nabla_\theta l^D_{\mathcal{T}_{tr}}(\theta_d)) \quad (9)$$

Therefore, the final update of the discriminator D for the batch is:

$$\theta_d \leftarrow \theta_d + \beta_1 \nabla_{\theta_d} \sum_{\mathcal{T}_{val} \in \mathcal{T}_i \sim p(\mathcal{T})} l^D_{\mathcal{T}_{val}}(\theta'_d) \quad (10)$$

Here, β_1 is the learning rate for the meta-step and θ'_d is the optimal parameter provided by the inner loop of meta-learner for the discriminator. Likewise, the generator's and classifier's objective function w.r.t. the unseen class examples in $\mathcal{T}_{val}$ is given by:

$$\min_{\theta_{gc}} \sum_{\mathcal{T}_{val} \in \mathcal{T}_i \sim p(\mathcal{T})} l^{GC}_{\mathcal{T}_{val}}(\theta'_{gc})$$
$$\xrightarrow{Update} \theta_{gc} \leftarrow \theta_{gc} - \beta_2 \nabla_{\theta_{gc}} \sum_{\mathcal{T}_{val} \in \mathcal{T}_i \sim p(\mathcal{T})} l^{GC}_{\mathcal{T}_{val}}(\theta'_{gc}) \quad (11)$$

Eq. 11 performs the meta-optimization across the batch of task for the generator and classifier. Again, note that, each task $\mathcal{T}_i = \{\mathcal{T}_{tr}, \mathcal{T}_{val}\}$ is partitioned into *training set* $\mathcal{T}_{tr}$ and *validation set* $\mathcal{T}_{val}$, such that the classes are disjoint. In contrast, traditional meta-learning (Finn, Abbeel, and Levine 2017) designed for few-shot learning assumes that the set of classes in $\mathcal{T}_{val}$ is same as the set of classes in $\mathcal{T}_{tr}$. This disjoint setup for $\mathcal{T}_{tr}$ and $\mathcal{T}_{val}$ is designed for zero-shot learning in order to mimic the problem setting which requires predicting the labels for examples from unseen classes not present at training time.

Example Generation and Zero-Shot Classification

After training the model, we can generate the unseen class examples given the respective class-attribute vectors. The generation of the novel class examples is done as:

$$\hat{\mathbf{x}} = G_{\theta_g}(\mathbf{z}, \mathbf{a}_c): \quad \mathbf{a}_c \in \mathbb{R}^d, c \in \{S+1, \ldots S+U\} \quad (12)$$

Here, $\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})$ and $\mathbf{z} \in \mathbb{R}^k$. Once we have generated samples from the unseen classes, we can train any classifier (e.g., SVM or softmax classifier) with these samples as labeled training data. In *generalized* ZSL setting, we synthesize samples from both seen and unseen class. We use the unseen class generated samples and actual/generated examples from seen classes to train a classifier with the label space being the union of seen and unseen classes. In practice, we found that using generated samples from seen classes (as opposed to actual samples) tends to perform better in the generalized ZSL setting. A justification for this is that the generated sample quality is uniform across seen and unseen class examples.

Related Work

Some of the earliest works on ZSL were based on directly or indirectly mapping the inputs to the class-attributes (Lampert, Nickisch, and Harmeling 2014; Norouzi et al. 2013; Socher et al. 2013). The learned mapping is used at inference time, this mapping first projects the unseen data to class-attribute space and then uses nearest neighbour search to predict the class. In a similar vein, other approaches (Romera and Torr 2015a; Changpinyo et al. 2016) also consider the relationship between seen and unseen classes. They represent the parameters of each unseen class as a similarity weighted combination of the parameters of seen classes. All of these models require plenty of data from the seen classes, and also do not work well in GZSL setting (Xian et al. 2018a).

Another prominent approach for ZSL focuses on learning the bilinear compatibility between the visual space and the semantic space of classes. (Akata et al. 2013; Frome et al. 2013; Akata et al. 2015; Romera and Torr 2015a; Kodirov, Xiang, and Gong 2017) are based on computing a linear/bilinear compatibility function. (Zhang and Saligrama 2015) embeds the inputs based on the semantic similarity. Some of the ZSL methods assume that all the unseen class inputs are also present at the time of training without the class labels. These *transductive* methods have extra information about all the unlabelled data of the unseen class, which leads to improved predictions as compared to the inductive setting (Song et al. 2018; Xu, Hospedales, and Gong 2017). Note that the transductive assumption is not very realistic since often test data is not available at the time of training.

The generalized ZSL (GZSL) (Verma et al. 2018; Chao et al. 2016; Xian et al. 2018a; 2018b) problem is arguably a very realistic and challenging problem wherein, unlike the ZSL problem, the training (seen) and the test (unseen) classes are not disjoint. Most of the previous models that perform well on standard ZSL fail to handle the biases towards predicting seen classes. Recently, generative models (Chen et al. 2018; Xian et al. 2018b; Verma and Rai 2017; Guo et

Method	SUN	CUB	AWA1	AWA2	aPY
LATEM (Xian et al. 2016)	55.3	49.3	55.1	55.8	35.2
ESZSL (Romera and Torr 2015a)	54.5	53.9	58.2	58.6	38.3
SYNC(Changpinyo et al. 2016)	56.3	55.6	54.0	46.6	23.9
DEM (Zhang, Xiang, and Gong 2017)	61.9	51.7	68.4	67.1	35.0
DCN (Liu et al. 2018)	61.8	56.2	–	65.2	43.6
ZSKL (Zhang and Koniusz 2018)	61.7	51.7	70.1	70.5	45.3
GFZSL(Verma and Rai 2017)	62.6	49.2	69.4	67.0	38.4
SP-AEN (Chen et al. 2018)	–	55.4	–	58.5	24.1
CVAE-ZSL(Mishra et al. 2017)	61.7	52.1	71.4	65.8	–
cycle-UWGAN (Felix et al. 2018)	59.9	58.6	–	66.8	–
f-CLSWGAN (Xian et al. 2018b)	60.8	57.3	–	68.2	–
SE-ZSL (Verma et al. 2018)	**63.4**	59.6	69.5	69.2	–
VSE-S (Zhu et al. 2019)	–	66.7	–	69.1	50.1
LisGAN (Li et al. 2019)	61.7	58.8	–	70.6	43.1
ZSML Softmax (Ours)	60.2	**69.6**	**73.5**	**76.1**	**64.1**
ZSML SVM (Ours)	60.1	**69.7**	**74.3**	**77.5**	**64.0**

Table 1: ZSL result using the per-class mean metric (Xian et al. 2018a). The non-generative models are mentioned at the top and the generative models are mentioned at the bottom. All compared methods use CNN-RNN feature for CUB dataset.

al. 2017; Wang et al. 2018) have shown promising results for both ZSL and GZSL setups. (Verma and Rai 2017) used a simple generative model based on the exponential family framework while (Guo et al. 2017) synthesized the classifier weights using class attributes. Recent generative approaches for ZSL are mostly based on VAE (Kingma and Welling 2014) and GAN (Goodfellow et al. 2014). Among these, (Verma et al. 2018; Bucher, Herbin, and Jurie 2017; Xian et al. 2019) are based on the VAE architectures while (Xian et al. 2018b; Chen et al. 2018; Li et al. 2019; Felix et al. 2018) use adversarial sample generation based on the class conditioned attribute. The recent approaches based on VAE and GAN show very competitive results. A particular advantage of the generative approaches is that, by using synthesized samples, we can convert the ZSL problem to the conventional supervised learning problem that can handle the biases towards the seen classes. The meta-learning approach are already tried for the ZSL (Hu, Xiong, and Socher 2018) to correct the learned network. To the best of our knowledge MAML (Finn, Abbeel, and Levine 2017) based approach over GAN has not been investigated yet. The meta-learning based adversarial generation model shows significant performance improvement, whereas the recent generative ZSL models have saturated.

Experiments and Results

We perform a comprehensive evaluation of our approach **ZSML** (**Z**ero-**S**hot **M**eta-**L**earning) by applying it on both standard ZSL and generalized ZSL problems and compare it with several state-of-the-art methods. We also perform several ablation studies to demonstrate/disentangle the benefits of the various aspects of our proposed approach. We evaluate our approach on the following benchmark ZSL datasets: SUN (Xiao et al. 2010) and CUB (Welinder et al. 2010) which are fine-grained and considered very challenging; AWA1 (Lampert, Nickisch, and Harmeling 2009) and AWA2 (Xian et al. 2018a); aPY (Farhadi et al. 2009) with diverse classes that makes this dataset very challenging. For CUB dataset, we use CNN-RNN textual features (Reed et

al. 2016) as class attributes, similar to the approaches mentioned in Table 1 and 2. Due to the lack of space, the complete *Algorithm* and details about the datasets are provided in the *Supplementary Material*. The generator and discriminator are 2-hidden layer networks with hidden layer size 2048 and 512, respectively.

Zero-Shot Learning

For the ZSL setting, we first train our model on seen class examples $\mathcal{D}^S$ and then synthesize samples from the unseen classes. These synthesized samples are further used to train either a multi-class linear SVM or a softmax classifier. The trained model over the synthesized examples is used to predict the classes for the test examples $\mathcal{D}^U$. We report results with both softmax classifier and linear SVM but we can, in principle, use any supervised classifier to train the model once we have generated the data. The average per-class accuracy is used as the standard evaluation metric (Xian et al. 2018a), shown in Table 1, as it overcomes the biases towards some particular class that has more data. In the ZSL setting, our model yields 4.5%, 6.0%, 9.8%, and 27.9% relative improvements over the current state-of-the-art on CUB, AWA1, AWA2, and aPY datasets, respectively. While, on the SUN dataset, it is very competitive as compared to the previous state-of-the-art methods. The SUN dataset contains 717 fine-grain classes; therefore, using the GAN based generation is highly prone to mode collapse. We believe that mode collapse is the possible reason for lower performance on SUN dataset. We are using the same network architecture and hyper-parameter for all the dataset. Since SUN dataset is fairly different compare to the other datasets, we believe that better hyper-parameter tuning for SUN dataset may improve the result.

Generalized Zero-Shot Learning

Standard ZSL assumes that all test inputs are from the unseen classes. The more challenging *generalized* Zero-Shot Learning (GZSL) relaxes this assumption and requires performing classification where the test set can potentially con-

Method	AWA1			CUB			aPY			AWA2		
	U	S	H	U	S	H	U	S	H	U	S	H
ESZSL (Romera and Torr 2015a)	6.6	75.6	12.1	12.6	63.8	21.0	2.4	70.1	4.6	5.9	77.8	11.0
SYNC (Changpinyo et al. 2016)	8.9	87.3	16.2	11.5	70.9	19.8	7.4	66.3	13.3	10.0	90.5	18.0
LATEM (Xian et al. 2016)	7.3	71.7	13.3	15.2	57.3	24.0	0.1	73.0	0.2	11.5	77.3	20.0
DEVISE (Frome et al. 2013)	13.4	68.7	22.4	23.8	53.0	32.8	4.9	76.9	9.2	17.1	74.7	27.8
DEM (Zhang, Xiang, and Gong 2017)	32.8	84.7	47.3	19.6	57.9	29.2	11.1	75.1	19.4	30.5	86.4	45.1
ZSKL (Zhang and Koniusz 2018)	18.3	79.3	29.8	21.6	52.8	30.6	10.5	76.2	18.5	18.9	82.7	30.8
DCN (Liu et al. 2018)	–	–	–	28.4	60.7	38.7	14.2	75.0	23.9	25.5	84.2	39.1
f-CLSWGAN (Xian et al. 2018b)	61.4	57.9	59.6	43.7	57.7	49.7	–	–	–	57.9	61.4	59.6
SP-AEN (Chen et al. 2018)	–	–	–	34.7	70.6	46.6	13.7	63.4	22.6	23.3	90.9	37.1
cycle-UWGAN (Felix et al. 2018)	–	–	–	47.9	59.3	53.0	–	–	–	59.6	63.4	59.8
SE-GZSL (Verma et al. 2018)	56.3	67.8	61.5	41.5	53.3	46.7	–	–	–	58.3	68.1	62.8
F-VAEGAND2 (Xian et al. 2019)	–	–	–	48.4	60.1	53.6	–	–	–	57.6	70.6	63.5
VSE-S (Zhu et al. 2019)	–	–	–	33.4	87.5	48.4	24.5	72.0	36.6	41.6	91.3	57.2
ZSML Softmax (Ours)	57.4	71.1	**63.5**	60.0	52.1	**55.7**	36.3	46.6	**40.9**	58.9	74.6	**65.8**

Table 2: Accuracy for GZSL, on novel proposed split (PS). U and S represent top-1 accuracy on unseen and seen class with all the $S + U$ classes. H stands for the harmonic mean. All compared methods use CNN-RNN feature for CUB dataset.

Method	N	AwA2			CUB		
		U	S	H	U	S	H
cycle-UWGAN	5	40.4	43.3	41.8	22.6	40.5	29.0
	10	45.5	50.9	48.0	25.5	42.1	32.5
f-CLSWGAN	5	37.8	44.2	40.7	30.4	28.5	29.4
	10	40.5	55.9	46.9	34.7	38.9	36.6
SE-GZSL	5	38.2	44.3	41.0	29.4	33.0	31.0
	10	41.4	45.1	43.1	35.6	43.5	39.1
Ours (ZSML)	5	44.5	54.8	**49.1**	31.4	38.1	**34.5**
	10	44.1	59.5	**50.7**	42.3	46.1	**44.1**

Table 3: GZSL results using only five and ten example per seen classes to train the model. The reported result is mean of 10 random split.

tain classes from the seen classes along with the unseen classes. We used the harmonic (HM) mean of the seen and unseen, average per class accuracy as the evaluation metric to report the results. It is found that HM (Xian et al. 2018a) is a better evaluation metric for GZSL since it overcomes the biases of predictions towards the seen class.

For GZSL task, we evaluate our model on the popular benchmark datasets CUB, aPY, AWA1 and AWA2. The results for GZSL is shown in Table 2. Our results demonstrate that ZSML achieves significant improvements in the harmonic mean. In terms of HM based accuracies, our ZSML yields 3.9%, 11.8%, 3.3% and 3.6% relative improvement over the current state-of-the-art on CUB, aPY, AWA1 and AWA2 datasets, respectively. Thus, ZSML not only works well in the standard ZSL setting but also in the GZSL setting. From Table 1 and 2, it is clear that all the models that show good results on the ZSL setup fail badly on the GZSL setup, whereas our model ZSML has consistently strong performance in both settings.

Ablation Study

In this section, we perform various ablation studies to assess the different aspects of our ZSML model on CUB, aPY and AWA2 datasets. We find that the proposed zero-shot meta-learning protocol (i.e., how we split the data from each task into meta-train and meta-validations sets) and meta-learning based adversarial generation are the key contributors for improving the model performance. We also conduct experiments when only few examples (say 5 or 10) are available from the seen class.

Meta-learner vs Plain-learner: We found that meta-learning based training is the key component to boost the model performance. Meta-learned model in the adversarial setting generates high-quality samples that are close to the real samples. In Figure 3, we are comparing the results with a recent approach (Chen et al. 2018; Felix et al. 2018; Xian et al. 2018b) that uses Improved-WGAN (Gulrajani et al. 2017) for the same problem.

To show the effectiveness of the proposed model, we are not using any advanced GAN architecture. We simply rely on the WGAN architecture. In the proposed model, the plain WGAN is associated with *meta-learning* agents. We have found that *meta-learning* framework is the key component to improve the performance. The proposed meta-learning framework improved the results in the ZSL setup, from 59.1% to 69.7% and 68.2% to 77.5% on CUB and AWA2 dataset respectively, compared to the current state-of-the-art as shown in Figure 3 (Top). Also in the same setting, our approach *without meta-learning* shows the ZSL results of 68.1% and 59.1% on AWA2 and CUB dataset respectively.

Few-Shot ZSL and Few-Shot GZSL: The meta-learning framework is specially designed for few-shot learning. So it is natural to ask how ZSL/GZSL will perform when only few-shot are present from the seen classes. This is the *most extreme* case for any classification algorithm (i.e. only a few examples are present from the seen class and at test time we have unseen/novel data). We perform the experiment for AWA2, CUB and aPY datasets assuming that only 5 or 10 examples per seen class are available. In the 5 examples per class experiment, we create a new dataset (by sampling from the original dataset) that contains 5 examples per seen classes (i.e. for 40 unseen classes in AWA2 dataset, our new

dataset contains only $5 \times 40 = 200$ samples). The model learns to generate unseen samples when it sees only 5 examples per seen class. Once the model is trained, we perform the classification following the procedure mentioned in Subsection . We follow the same process for 10 examples per seen class. As shown in Figure 3 (Bottom), with as few as only 10 examples per-class our approach outperforms other state-of-the-art methods on CUB, aPY and AWA2 datasets in ZSL setting, also using only 5 examples per class our result are very competitive (while competitor model uses *all examples* in training). Also as shown in Table 3, in the most challenging GZSL setting, using only 5 or 10 samples our result out performs the recent approach by a significant margin.

Zero-Shot MAML Split vs Traditional MAML Split: We propose a novel task distribution for ZSML where each task $\mathcal{T}_i$ is partitioned into two sets $\mathcal{T}_{tr}$ and $\mathcal{T}_{val}$ and the classes in $\mathcal{T}_{tr}$ and $\mathcal{T}_{val}$ are disjoint. While in the MAML setup these classes are the same. The ablation over the MAML and ZSML task distribution is shown in Figure 2. The proposed training set and validation set split (per episode) performs significantly better than traditional MAML split. Using the novel ZSML split, the ZSL results improves 1.7% and 2.4% on the AWA2 and CUB dataset, respectively.

Which Aspects Benefit More from Meta-Adversarial Learning? In adversarial learning, the sample quality depends on how powerful the discriminator and generator are. The optimal discriminator minimizes the JS-Divergence between the generated and the original samples (Goodfellow et al. 2014). The meta-learner associated with discriminator or generator provides a powerful discriminator and generator by enhancing their learning capability. The optimal discriminator provides strong feedback to the generator and the generator continuously increases its generation capability. We observe that if we remove the meta-learner from the discriminator, we have 5.8% and 8.6% accuracy drop as compared to our model with a meta-learning component on CUB and AWA2 dataset, respectively. The significant accuracy drop occurs since the discriminator is not optimal and provides poor feedback to the generator. Similarly, if we remove the meta-learner from the generator, we again observe a significant accuracy drop (2.2% and 7.9% on CUB and AWA2 dataset, respectively). Since the generator has a reduced capability without meta-learner, even though discriminator provides strong feedback to the generator, the gen-


Figure 2: Our ZSL result for AWA2 and CUB datasets with the proposed zero-shot task distribution.
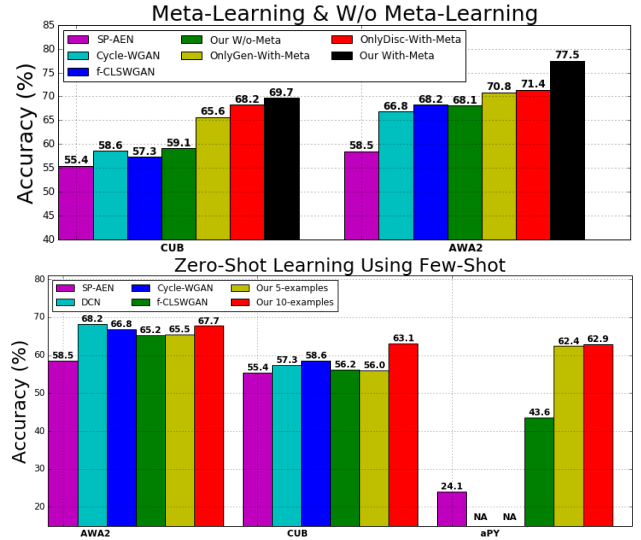

Figure 3: **Top:** Comparison of ZSL results on AWA2 and CUB dataset with recently proposed models based on GAN and our meta-learned GAN. **Bottom:** Our ZSL result using only few samples (say 5 and 10) per seen class, compared to all other methods that use all the samples.

erator is not powerful enough to counter the discriminator. Also, if we remove the meta-learning agent from generator and discriminator, it becomes a plain adversarial network. The ablation results are shown in Figure 3.

Conclusion

In this work, we identify and address three key limitations of current ZSL approaches, that limit the performance of the recent generative models for ZSL/GZSL. We observe that a meta-learning based approach can naturally overcome these limitations in a principled manner. We have proposed a novel framework for ZSL and GZSL which is based on the meta-learning framework over a conditional generative model (WGAN). We also propose a novel zero-shot task distribution for the meta-learning model to mimic the ZSL behaviour. We have conducted extensive experiments benchmark ZSL datasets. In the few-shot, as well as standard GZSL setting, the proposed model outperforms the state-of-the-art methods by a significant margin. Our ablation study shows that the proposed meta-learning framework and zero-shot task distribution are the key components for performance improvement. Finally, although our focus here has been on ZSL and generalized ZSL, our meta-learning based adversarial generation model can be useful for the problem of distribution learning and generation tasks as well (Hewitt et al. 2018). For GZSL, we achieve the state-of-the-art results over all the standard datasets, whereas for ZSL, we surpass the state-of-art results by a significant margin on aPY, CUB, AWA1 and AWA2 datasets.

References

Akata, Z.; Perronnin, F.; Harchaoui, Z.; and Schmid, C. 2013. Label-embedding for attribute-based classification. In *CVPR*, 819.

Akata, Z.; Reed, S.; Walter, D.; Lee, H.; and Schiele. 2015. Evaluation of output embeddings for fine-grained image classification. In *CVPR*, 2927–2936.

Arjovsky, M.; Chintala, S.; and Bottou, L. 2017. Wasserstein gan. *arXiv preprint arXiv:1701.07875*.

Bucher, M.; Herbin, S.; and Jurie, F. 2017. Generating visual representations for zero-shot classification. In *ICCV, Workshop*.

Changpinyo, S.; Chao, W.-L.; Gong, B.; and Sha, F. 2016. Synthesized classifiers for zero-shot learning. In *CVPR*, 5327–5336.

Chao, W.-L.; Changpinyo, S.; Gong, B.; and Sha, F. 2016. An empirical study and analysis of generalized zero-shot learning for object recognition in the wild. In *ECCV*.

Chen, L.; Zhang, H.; Xiao, J.; Liu, W.; and Chang, S.-F. 2018. Zero-shot visual recognition using semantics-preserving adversarial embedding networks. In *CVPR*.

Clouâtre, L., and Demers, M. 2019. Figr: Few-shot image generation with reptile. *arXiv preprint arXiv:1901.02199*.

Farhadi, A.; Endres, I.; Hoiem, D.; and Forsyth, D. 2009. Describing objects by their attributes. In *CVPR*, 1778–1785. IEEE.

Felix, R.; Vijay Kumar, B.; Reid, I.; and Carneiro, G. 2018. Multi-modal cycle-consistent generalized zero-shot learning. *ECCV*.

Finn, C.; Abbeel, P.; and Levine, S. 2017. Model-agnostic meta-learning for fast adaptation of deep networks. *ICML*.

Frome, A.; Corrado, G. S.; Shlens, J.; Bengio, S.; Dean, J.; Mikolov, T.; et al. 2013. Devise: A deep visual-semantic embedding model. In *NIPS*, 2121–2129.

Goodfellow, I.; Pouget-Abadie, J.; Mirza, M.; Xu, B.; Warde-Farley, D.; Ozair, S.; Courville, A.; and Bengio, Y. 2014. Generative adversarial nets. In *NIPS*, 2672–2680.

Gulrajani, I.; Ahmed, F.; Arjovsky, M.; Dumoulin, V.; and Courville, A. C. 2017. Improved training of wasserstein gans. In *NIPS*, 5767–5777.

Guo, Y.; Ding, G.; Han, J.; and Gao, Y. 2017. Synthesizing samples for zero-shot learning. In *IJCAI*.

Hewitt, L. B.; Nye, M. I.; Gane, A.; Jaakkola, T.; and Tenenbaum, J. B. 2018. The variational homoencoder: Learning to learn high capacity generative models from few examples. *arXiv preprint arXiv:1807.08919*.

Hu, R. L.; Xiong, C.; and Socher, R. 2018. Correction networks: Meta-learning for zero-shot learning.

Kingma, D. P., and Welling, M. 2014. Auto-encoding variational bayes. *ICLR*.

Kodirov, E.; Xiang, T.; and Gong, S. 2017. Semantic autoencoder for zero-shot learning. *CVPR*.

Lampert, C. H.; Nickisch, H.; and Harmeling, S. 2009. Learning to detect unseen object classes by between-class attribute transfer. In *CVPR*, 951–958. IEEE.

Lampert, C. H.; Nickisch, H.; and Harmeling, S. 2014. Attribute-based classification for zero-shot visual object categorization. *PAMI* 36(3):453–465.

Li, J.; Jin, M.; Lu, K.; Ding, Z.; Zhu, L.; and Huang, Z. 2019. Leveraging the invariant side of generative zero-shot learning. *CVPR*.

Liu, S.; Long, M.; Wang; and Jordan, M. I. 2018. Generalized zero shot learning with deep calibration network. In *NIPS*, 2006–16.

Mishra, A.; Reddy, M.; Mittal, A.; and Murthy, H. A. 2017. A generative model for zero shot learning using conditional variational autoencoders. *CVPR Workshop*.

Norouzi, M.; Mikolov, T.; Bengio, S.; Singer, Y.; Shlens, J.; Frome, A.; Corrado, G. S.; and Dean, J. 2013. Zero-shot learning by convex combination of semantic embeddings. *NIPS*.

Reed, S.; Akata, Z.; Lee, H.; and S, B. 2016. Learning deep representations of fine-grained visual descriptions. In *CVPR*, 49–58.

Romera, Paredes, B., and Torr, P. 2015a. An embarrassingly simple approach to zero-shot learning. In *ICML*, 2152–2161.

Romera, Paredes, B., and Torr, P. H. 2015b. An embarrassingly simple approach to zero-shot learning. In *ICML*, 2152–2161.

Snell, J.; Swersky, K.; and Zemel, R. 2017. Prototypical networks for few-shot learning. In *NIPS*, 4077–4087.

Socher, R.; Ganjoo, M.; Manning, C. D.; and Ng, A. 2013. Zero-shot learning through cross-modal transfer. In *NIPS*, 935–943.

Song, J.; Shen, C.; Yang, Y.; Liu, Y.; and S, M. 2018. Transductive unbiased embedding for zero-shot learning. In *CVPR*, 1024–1033.

Verma, V. K., and Rai, P. 2017. A simple exponential family framework for zero-shot learning. In *ECML-PKDD*, 792–808. Springer.

Verma, V. K.; Arora, G.; Mishra, A.; and Rai, P. 2018. Generalized zero-shot learning via synthesized examples. *CVPR*.

Vinyals, O.; Blundell, C.; Lillicrap, T.; Wierstra, D.; et al. 2016. Matching networks for one shot learning. In *NIPS*, 3630–3638.

Wang, W.; Pu, Y.; Verma, V. K.; Fan, K.; Zhang, Y.; Chen, C.; Rai, P.; and Carin, L. 2018. Zero-shot learning via class-conditioned deep generative models. *AAAI*.

Welinder, P.; Branson, S.; Mita, T.; Wah, C.; Schroff, F.; Belongie, S.; and Perona, P. 2010. Caltech-ucsd birds 200. *California Institute of Technology*.

Xian, Y.; Akata, Z.; Sharma, G.; Nguyen, Q.; Hein, M.; and Schiele, B. 2016. Latent embeddings for zero-shot classification. In *CVPR*, 69–77.

Xian, Y.; Lampert, C. H.; Schiele, B.; and Akata, Z. 2018a. Zero-shot learning-a comprehensive evaluation of the good, the bad and the ugly. *PAMI*.

Xian, Y.; Lorenz, T.; Schiele, B.; and Akata, Z. 2018b. Feature generating networks for zero-shot learning. In *CVPR*.

Xian, Y.; Sharma, S.; Schiele, B.; and Akata, Z. 2019. f-vaegan-d2: A feature generating framework for any-shot learning. In *CVPR*, 10275–10284.

Xiao, J.; Hays, J.; Ehinger, K. A.; Oliva, A.; and Torralba, A. 2010. Sun database: Large-scale scene recognition from abbey to zoo. In *CVPR, 2010*, 3485–3492.

Xu, X.; Hospedales, T.; and Gong, S. 2017. Transductive zero-shot action recognition by word-vector embedding. *International Journal of Computer Vision* 1–25.

Zhang, H., and Koniusz, P. 2018. Zero-shot kernel learning. In *CVPR*, 7670–7679.

Zhang, Z., and Saligrama, V. 2015. Zero-shot learning via semantic similarity embedding. In *ICCV*, 4166–4174.

Zhang, L.; Xiang, T.; and Gong, S. 2017. Learning a deep embedding model for zero-shot learning. In *CVPR*, 3010–3019.

Zhu, Y.; Elhoseiny, M.; Liu, B.; Peng, X.; and Elgammal, A. 2018. A generative adversarial approach for zero-shot learning from noisy texts. In *CVPR*, 1004–1013.

Zhu; Pengkai; Wang, H.; and Saligrama, V. 2019. Generalized zero-shot recognition based on visually semantic embedding. In *CVPR*, 2995–3003.