

SELF-TRAINING FOR FEW-SHOT TRANSFER ACROSS EXTREME TASK DIFFERENCES

Cheng Perng Phoo, Bharath Hariharan

Department of Computer Science

Cornell University

{cpphoo, bharathh}@cs.cornell.edu

ABSTRACT

All few-shot learning techniques must be pre-trained on a large, labeled “base dataset”. In problem domains where such large labeled datasets are not available for pre-training (e.g., X-ray images), one must resort to pre-training in a different “source” problem domain (e.g., ImageNet), which can be very different from the desired target task. Traditional few-shot and transfer learning techniques fail in the presence of such extreme differences between the source and target tasks. In this paper, we present a simple and effective solution to tackle this extreme domain gap: self-training a source domain representation on unlabeled data from the target domain. We show that this improves one-shot performance on the target domain by 2.9 points on average on a challenging benchmark with multiple domains.

1 INTRODUCTION

Despite progress in visual recognition, training recognition systems for new classes in novel domains requires thousands of labeled training images per class and several hours of compute. For example, to train a recognition system for different kinds of pneumonia in chest X-rays, one would have to get radiologists to label thousands of X-ray images, and then spend several hours to train a neural network on high-end GPUs. The high cost of doing so precludes many downstream applications.

This issue has motivated research on *few-shot learners*: systems that can *rapidly* learn novel classes from *a few examples*. However, all few-shot learners must be trained on a large *base dataset* of classes from the same domain. This is a problem in many domains (such as medical imagery, satellite images), where no large labeled dataset of base classes exists. The only alternative is to train the few-shot learner on a different domain (a common choice is to use ImageNet). Unfortunately, *few-shot learning techniques assume that novel and base classes share modes of variation* (Wang et al., 2018), *class-distinctive features* (Snell et al., 2017), *or other inductive biases*. These assumptions are broken when the difference between base and novel is as extreme as the difference between object classification in internet photos and pneumonia detection in X-ray images. As such, recent work has found that all few-shot learners fail in the face of such extreme task/domain differences, underperforming even naive transfer learning from ImageNet (Guo et al., 2020).

Another alternative comes to light when one considers that many of these problem domains have *unlabeled data* (e.g., undiagnosed X-ray images, or unlabeled satellite imagery). This suggests the possibility of using *self-supervised techniques* on this unlabeled data to produce a good feature representation, which can then be used to train linear classifiers for the target classification task using just a few labeled examples. Indeed, recent work has explored self-supervised learning on a variety of domains (Wallace & Hariharan, 2020). However, self-supervised learning starts *tabula rasa*, and as such requires extremely large amounts of unlabeled data (on the order of millions of images). With more practical unlabeled datasets, self supervised techniques still struggle to outcompete naive ImageNet transfer (Wallace & Hariharan, 2020). We are thus faced with a conundrum: on the one hand, few-shot learning techniques fail to bridge the extreme differences between ImageNet and domains such as X-rays. On the other hand, self-supervised techniques fail when they ignore inductive biases from ImageNet. A sweet spot in the middle, if it exists, is elusive.

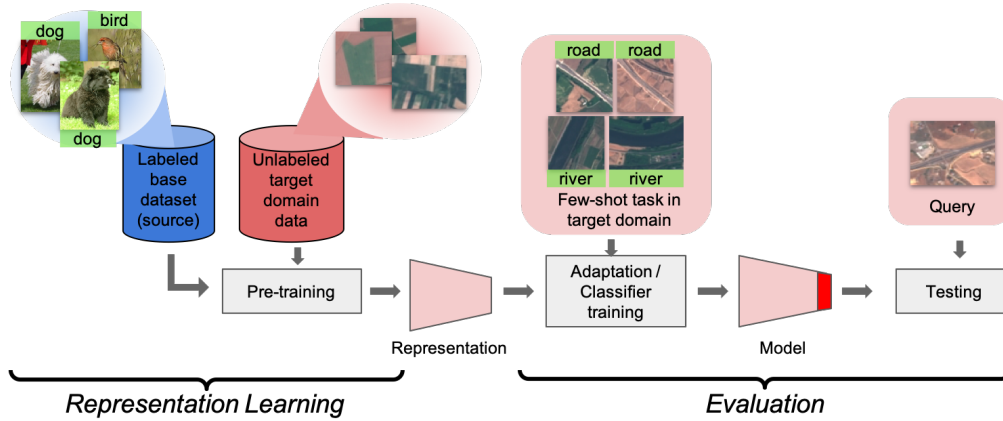


Figure 1: Problem setup. In the representation learning phase (left), the learner has access to a large labeled “base dataset” in the source domain, and some unlabeled data in the target domain, on which to pre-train its representation. The learner must then rapidly learn/adapt to few-shot tasks in the target domain in the evaluation phase (right).

In this paper, we solve this conundrum by presenting a strategy that adapts feature representations trained on source tasks to *extremely different* target domains, so that target task classifiers can then be trained on the adapted representation with very little labeled data. Our key insight is that a pre-trained base classifier from the source domain, when applied to the target domain, induces a grouping of images on the target domain. This grouping captures what the pre-trained classifier thinks are similar or dissimilar in the target domain. Even though the classes of the pre-trained classifier are themselves irrelevant in the target domain, the induced notions of *similarity and dissimilarity* might still be relevant and informative. This induced notion of similarity is in contrast to current self-supervised techniques which often function by considering each image as its own class and dissimilar from every other image in the dataset (Wu et al., 2018; Chen et al., 2020). We propose to train feature representations on the novel target domain to *replicate this induced grouping*. This approach produces a feature representation that is (a) adapted to the target domain, while (b) maintaining prior knowledge from the source task to the extent that it is relevant. A discerning reader might observe the similarity of this approach to *self-training*, except that our goal is to adapt the feature representation to the target domain, rather than improve the base classifier itself.

We call our approach “Self Training to Adapt Representations To Unseen Problems”, or STARTUP. In a recently released benchmark consisting of datasets from extremely different domains (Guo et al., 2020), we show that STARTUP provides significant gains (up to 2.9 points on average) over few-shot learning, transfer learning and self-supervision state-of-the-art. To the best of our knowledge, ours is the first attempt to bridge such large task/domain gaps and successfully and consistently outperform naive transfer in cross-domain few-shot learning.

2 PROBLEM SETUP

Our goal is to build learners for novel domains that can be *quickly* trained to recognize new classes when presented with *very few* labeled data points (“*few-shot*”). Formally, the target domain is defined by a set of data points (e.g. images) $\mathcal{X}_\mathcal{N}$, an unknown set of classes (or label space) $\mathcal{Y}_\mathcal{N}$, and a distribution $\mathcal{D}_\mathcal{N}$ over $\mathcal{X}_\mathcal{N} \times \mathcal{Y}_\mathcal{N}$. A “few-shot learning task” in this domain will consist of a set of classes $Y \subset \mathcal{Y}_\mathcal{N}$, a very small training set (“support”)

$$S = \{(x_i, y_i)\}_{i=1}^n \sim \mathcal{D}_\mathcal{N}^n, \quad y_i \in Y$$

and a small test set (“query”)

$$Q = \{x_i\}_{i=1}^m \sim \mathcal{D}_\mathcal{N}^m$$

When presented with such a few-shot learning task, the learner must rapidly learn the classes presented and accurately classify the query images.

As with prior few-shot learning work, we will assume that before being presented with few-shot learning tasks in the target domain, the learner has access to a large annotated dataset D_B known as the base dataset. However, crucially *unlike prior work on few-shot learning*, we assume that this base dataset is drawn from a very different distribution. In fact, we assume that the base dataset is drawn from a completely disjoint image space \mathcal{X}_B and a disjoint set of classes \mathcal{Y}_B :

$$D_B = \{(x_i, y_i)\}_{i=1}^{N_B} \subset \mathcal{X}_B \times \mathcal{Y}_B$$

where \mathcal{X}_B is the set of data (or the source domain) and \mathcal{Y}_B is the set of base classes. Because the base dataset is so different from the target domain, we introduce another difference vis-a-vis the conventional few-shot learning setup: the learner is given access to an additional unlabeled dataset from the target domain:

$$D_u = \{x_i\}_{i=1}^{N_u} \sim \mathcal{D}_{\mathcal{N}}^{N_u}$$

Put together, the learner will undergo two phases. In the *representation learning* phase, the learner will pre-train its representation on D_B and D_u ; then it goes into the *evaluation* phase where it will be presented few-shot tasks from the target domain where it learns the novel classes (Figure 1).

3 RELATED WORK

Few-shot Learning (FSL). This paper explores few-shot transfer, and as such the closest related work is on few-shot learning. Few-shot learning techniques are typically predicated on some degree of similarity between classes in the base dataset and novel classes. For example, they may assume that features that are discriminative for the base classes are also discriminative for the novel classes, suggesting a metric learning-based approach (Gidaris & Komodakis, 2018; Qi et al., 2018; Snell et al., 2017; Vinyals et al., 2016; Sung et al., 2018; Hou et al., 2019). Alternatively, they may assume that model initializations that lead to rapid convergence on the base classes are also good initializations for the novel classes (Finn et al., 2017; 2018; Ravi & Larochelle, 2017; Nichol & Schulman; Rusu et al., 2019; Sun et al., 2019; Lee et al., 2019). Other methods assume that modes of intra-class variation are shared, suggesting the possibility of learned, class-agnostic augmentation policies (Hariharan & Girshick, 2017; Wang et al., 2018; Chen et al., 2019b). Somewhat related is the use of a class-agnostic parametric model that can “denoise” few-shot models, be they from the base or novel classes (Gidaris & Komodakis, 2018; 2019). In contrast to such strong assumptions of similarity between base and novel classes, this paper tackles few-shot learning problems where base and novel classes come from very different domains, also called cross-domain few-shot learning.

Cross-domain Few-shot Classification (CD-FSL). When the domain gap between the base and novel dataset is large, recent work (Guo et al., 2020; Chen et al., 2019a) has shown that existing state-of-the-art few-shot learners fail to generalize. Tseng et al. (2020) attempt to address this problem by simulating cross-domain transfer during training. However, their approach assumes access to an equally diverse array of domains during training, and a much smaller domain gap at test time: for example, both base and novel datasets are from internet images. Our paper tackles a more extreme domain gap.

Few-shot learning with unlabeled data. This paper uses unlabeled data from the target domain to bridge the domain gap. Semi-supervised few-shot learning (SS-FSL) (Ren et al., 2018; Li et al., 2019; Yu et al., 2020; Rodríguez et al., 2020; Wang et al., 2020) and transductive few-shot learning (T-FSL) (Liu et al., 2019; Dhillon et al., 2020; Hou et al., 2019; Wang et al., 2020; Rodríguez et al., 2020) do use such unlabeled data, but only during evaluation, assuming that representations trained on the base dataset are good enough. In contrast our approach leverages the unlabeled data during representation learning. The two are orthogonal innovations and can be combined.

Self-Training. Our approach is closely related to self-training, which has been shown to be effective for semi-supervised training and knowledge distillation. In self-training, a teacher model trained on the labeled data is used to label the unlabeled data and another student model is trained on both the original labeled data and the unlabeled data labeled by the teacher. Xie et al. (2020) and Yalniz et al. (2019) have shown that using self-training can improve ImageNet classification performance. Knowledge distillation (Hinton et al., 2015) is similar but aims to compress a large teacher network by training a student network to mimic the prediction of the teacher network. A key difference between these and our work is that self-training / knowledge distillation focus on a single task of

interest, i.e., there is no change in label space. Our approach is similar, but we are interested in transferring to novel domains with a *wholly different label space*: an unexplored scenario.

Domain Adaptation. Transfer to new domains is also in the purview of domain adaptation (Tzeng et al., 2017; Hoffman et al., 2018; Long et al., 2018; Xu et al., 2019; Laradji & Babanezhad, 2020; Wang & Deng, 2018; Wilson & Cook, 2020) where the goal is to transfer knowledge from the label-abundant source domain to a target domain where only unlabeled data is available. However, a key assumption in domain adaptation is that the source domain and target domain share the same label space which does not hold for few-shot learning.

Self-supervised Learning. Learning from unlabeled data has seen a resurgence of interest with advances in self-supervised learning. Early self-supervised approaches were based on handcrafted “pretext tasks” such as solving jigsaw puzzles (Noroozi & Favaro, 2016), colorization (Zhang et al., 2016) or predicting rotation (Gidaris et al., 2018). A more recent (and better performing) line of self-supervised learning is contrastive learning (Wu et al., 2018; Misra & Maaten, 2020; He et al., 2020; Chen et al., 2020) which aims to learn representations by considering each image together with its augmentations as a separate class. While self supervision has been shown to boost few-shot learning (Gidaris et al., 2019; Su et al., 2020), its utility in cases of large domain gaps between base and novel datasets have not been evaluated. Our work focuses on this challenging scenario.

4 APPROACH

Consider a classification model $f_\theta = C \circ \phi$ where ϕ embeds input x into \mathbb{R}^d and C is a (typically linear) classifier head that maps $\phi(x)$ to predicted probabilities $P(y|x)$. θ is a vector of parameters. During representation learning, STARTUP performs the following three steps:

1. Learn a teacher model θ_0 on the base dataset \mathcal{D}_B by minimizing the cross entropy loss
2. Use the teacher model to construct a **softly-labeled set** $\mathcal{D}_u^* = \{(x_i, \bar{y}_i)\}_{i=1}^{N_u}$ where

$$\bar{y}_i = f_{\theta_0}(x_i) \quad \forall x_i \in \mathcal{D}_u. \quad (1)$$

Note that \bar{y}_i is a probability distribution as described above.

3. Learn a new student model θ^* on \mathcal{D}_B and \mathcal{D}_u^* by optimizing:

$$\min_{\theta} \frac{1}{N_B} \sum_{(x_i, y_i) \in \mathcal{D}_B} l_{CE}(f_\theta(x_i), y_i) + \frac{1}{N_u} \sum_{(x_j, \bar{y}_j) \in \mathcal{D}_u^*} l_{KL}(f_\theta(x_j), \bar{y}_j) + l_{unlabeled}(\mathcal{D}_u) \quad (2)$$

where l_{CE} is the cross entropy loss, l_{KL} is the KL divergence and $l_{unlabeled}$ is any unsupervised/self-supervised loss function (See below).

The third term, $l_{unlabeled}$, is intended to help the learner extract additional useful knowledge specific to the target domain. We use a state-of-the-art self-supervised loss function based on contrastive learning: *SimCLR* (Chen et al., 2020). The SimCLR loss encourages two augmentations of the same image to be closer in feature space to each other than to other images in the batch. We refer the reader to the paper for the detailed loss formulation.

The first two terms are similar to those in prior self-training literature (Xie et al., 2020). However, while in prior self-training work, the second term (l_{KL}) is thought to mainly introduce noise during training, we posit that l_{KL} has a more substantial role to play here: it encourages the model to learn feature representations that *emphasize* the groupings induced by the pseudo-labels \bar{y}_i on the target domain. We analyze this intuition in section 5.2.2.

4.1 EVALUATION

STARTUP is agnostic to inference methods during evaluation; any inference methods that rely on a representation (Snell et al., 2017; Gidaris & Komodakis, 2018) can be used with STARTUP. For simplicity and based on results reported by Guo et al. (2020), we freeze the representation ϕ after performing STARTUP and train a linear classifier on the support set and evaluate the classifier on the query set.

4.2 INITIALIZATION STRATEGIES

Xie et al. (2020) found that training the student from scratch sometimes yields better results for ImageNet classification. For STARTUP, the student is initialized to the teacher embedding with a randomly initialized classifier by default. We experimented with two other initialization strategies - from scratch (STARTUP-Rand) and from teacher (STARTUP-T) and find (STARTUP-Rand) hurts performance whereas STARTUP-T performs slightly worse than STARTUP (See Appendix A.3).

5 EXPERIMENTS

We defer the implementation details to Appendix A.1.

5.1 FEW-SHOT TRANSFER ACROSS DRASTICALLY DIFFERENT DOMAINS

Benchmark. We experiment with the challenging (BSCD-FSL) benchmark introduced in Guo et al. (2020). The *base dataset* in this benchmark is miniImageNet (Vinyals et al., 2016), which is an object recognition task on internet images. There are 4 novel datasets in the benchmark, none of which involve objects, and all of which come from a very different domain than internet images: CropDiseases (recognizing plant diseases in leaf images), EuroSAT (predicting land-use from satellite images), ISIC2018 (identifying melanoma from images of skin lesions) and ChestX (diagnosing chest X-rays). Guo et al. found that state-of-the-art few-shot learners fail on this benchmark.

To construct our setup, we randomly sample 20% of data from each novel datasets to form the respective unlabeled datasets \mathcal{D}_u . We use the rest for sampling tasks for evaluation. Following Guo et al. (2020), we evaluate 5-way k-shot classification tasks (the support set consists of 5 classes and k examples per class) for $k \in \{1, 5\}$ and report the mean and 95% confidence interval over 600 few-shot tasks (conclusions generalize to $k \in \{20, 50\}$. See Appendix A.2).

Baselines. We compare to the techniques reported in Guo et al. (2020), which includes most state-of-the-art approaches as well as a cross-domain few-shot technique Tseng et al. (2020). The top performing among these is naive **Transfer** which simply trains a convolutional network to classify the base dataset, and uses the resulting representation to learn a linear classifier when faced with novel few-shot tasks. These techniques do not use the novel domain unlabeled data.

We also compare to an additional baseline, **SimCLR** that uses the novel domain unlabeled data \mathcal{D}_u to train a representation using SimCLR(Chen et al., 2020), and then uses the resulting representation to learn linear classifiers for few-shot tasks. This builds upon state-of-the-art self-supervised techniques.

Following the benchmark, all methods use a ResNet-10 (He et al., 2016) unless otherwise stated.

5.1.1 RESULTS

We present our main results on miniImageNet \rightarrow BSCD-FSL in Table 1.

STARTUP vs Few-shot learning techniques. STARTUP performs significantly better than all few-shot techniques in most datasets (except ChestX, where all methods are similar). Compared to previous state-of-the-art, **Transfer**, we observe an average of 2.9 points improvement on the 1-shot case. The improvement is particularly large on CropDisease, where STARTUP provides almost a 6 point increase for 1-shot classification. This improvement is significant given the simplicity of our approach, and given that all meta-learning techniques *underperform* this baseline.

STARTUP vs SimCLR. The SimCLR baseline in general tends to *underperform* naive transfer from miniImageNet, and consequently, STARTUP performs significantly better than SimCLR on ISIC and EuroSAT. The exception to this is CropDisease, where SimCLR produces a surprisingly good representation. We conjecture that the base embedding is not a good starting point for this dataset.

Different variants of STARTUP. In general all variants of STARTUP outperform the transfer baseline. However, we find that using SimCLR as an auxiliary loss to train the student (STARTUP vs STARTUP (no SS), and STARTUP-T vs STARTUP-T (no SS)) is beneficial. Additionally,

Table 1: 5-way k-shot classification accuracy on miniImageNet→BSCD-FSL. Mean and 95% confidence interval are reported. (no SS) indicates removal of SimCLR. ProtoNet: (Snell et al., 2017), MAML:(Finn et al., 2017), MetaOpt:(Lee et al., 2019) FWT:(Tseng et al., 2020). *Numbers reported in (Guo et al., 2020); our re-implementation of Transfer uses a different batch size and 80% of the original test set for evaluation.

Methods	ChestX		ISIC	
	k=1	k=5	k=1	k=5
MAML*	-	23.48 \pm 0.96	-	40.13 \pm 0.58
ProtoNet*	-	24.05 \pm 1.01	-	39.57 \pm 0.57
ProtoNet + FWT*	-	23.77 \pm 0.42	-	38.87 \pm 0.52
MetaOpt*	-	22.53 \pm 0.91	-	36.28 \pm 0.50
Transfer*	-	25.35 \pm 0.96	-	43.56 \pm 0.60
Transfer	22.71 \pm 0.40	26.71 \pm 0.46	30.71 \pm 0.59	43.08 \pm 0.57
SimCLR	22.10 \pm 0.41	25.02 \pm 0.42	26.25 \pm 0.53	36.09 \pm 0.57
STARTUP-T (no SS)	22.79 \pm 0.41	26.03 \pm 0.43	32.37 \pm 0.61	45.20 \pm 0.61
STARTUP (no SS)	22.87 \pm 0.41	26.68 \pm 0.45	32.24 \pm 0.62	46.48 \pm 0.61
STARTUP-T	22.75 \pm 0.40	26.47 \pm 0.43	32.16 \pm 0.60	45.75 \pm 0.60
STARTUP	23.09 \pm 0.43	26.94 \pm 0.44	32.66 \pm 0.60	47.22 \pm 0.61

Methods	EuroSAT		CropDisease	
	k=1	k=5	k=1	k=5
MAML*	-	71.70 \pm 0.72	-	78.05 \pm 0.68
ProtoNet*	-	73.29 \pm 0.71	-	79.72 \pm 0.67
ProtoNet + FWT*	-	67.34 \pm 0.76	-	72.72 \pm 0.70
MetaOpt*	-	64.44 \pm 0.73	-	68.41 \pm 0.73
Transfer*	-	75.69 \pm 0.66	-	87.48 \pm 0.58
Transfer	60.73 \pm 0.86	80.30 \pm 0.64	69.97 \pm 0.85	90.16 \pm 0.49
SimCLR	43.52 \pm 0.88	59.05 \pm 0.70	78.23 \pm 0.83	92.57 \pm 0.48
STARTUP-T (no SS)	63.00 \pm 0.84	81.25 \pm 0.62	71.11 \pm 0.83	90.79 \pm 0.49
STARTUP (no SS)	62.90 \pm 0.83	81.81 \pm 0.61	73.30 \pm 0.82	91.69 \pm 0.47
STARTUP-T	63.49 \pm 0.85	81.54 \pm 0.63	72.85 \pm 0.83	91.49 \pm 0.48
STARTUP	63.88 \pm 0.84	82.29 \pm 0.60	75.93 \pm 0.80	93.02 \pm 0.45

Table 2: 5-way k-shot classification accuracy on ImageNet(ILSVRC 2012)→BSCD-FSL.

Methods	ChestX		ISIC	
	k=1	k=5	k=1	k=5
Transfer	21.97 \pm 0.39	25.85 \pm 0.41	30.27 \pm 0.51	43.88 \pm 0.56
STARTUP (no SS)	22.90 \pm 0.40	26.74 \pm 0.46	30.18 \pm 0.56	44.19 \pm 0.57
STARTUP	23.03 \pm 0.42	27.24 \pm 0.46	31.69 \pm 0.59	46.02 \pm 0.59

Methods	EuroSAT		CropDisease	
	k=1	k=5	k=1	k=5
Transfer	66.08 \pm 0.81	85.58 \pm 0.48	74.17 \pm 0.82	92.46 \pm 0.42
STARTUP (no SS)	70.08 \pm 0.80	87.12 \pm 0.45	80.13 \pm 0.77	94.51 \pm 0.38
STARTUP	73.83 \pm 0.77	89.70 \pm 0.41	85.10 \pm 0.74	96.06 \pm 0.33

Table 3: Adjusted Mutual Information (AMI) of the grouping induced by the teacher and the ground truth label. AMI has value from 0 to 1 with higher value indicating more agreement.

	ChestX	ISIC	EuroSAT	CropDisease
AMI	0.0075	0.0427	0.3079	0.2969

re-initializing the classifier head works better than using the pre-trained classifier (STARTUP vs STARTUP-T), potentially because the classifier is too overfitted to the base dataset.

Larger and stronger teachers. To unpack the impact of teacher quality, we experiment with a larger network and transfer from the full ILSVRC 2012 dataset (Deng et al., 2009) to BSCD-FSL. In particular, we used the publicly available pre-trained ResNet-18 (He et al., 2016) as a teacher and train a student via STARTUP. We compare this to a transfer baseline that uses the same network and ImageNet as the training set. The result can be found in table 2. Surprisingly, larger, richer embeddings *do not always transfer better*, in contrast to in-domain results reported by Hariharan & Girshick (2017). However, STARTUP is still useful in improving performance: the absolute improvement in performance for STARTUP compared to Transfer remains about the same in most datasets except EuroSAT and CropDisease where larger improvements are observed.

5.2 WHY SHOULD STARTUP WORK?

While it is clear that STARTUP helps improve few shot transfer across extreme domain differences, it is not clear why or how it achieves this improvement. Below, we look at a few possible hypotheses.

5.2.1 HYPOTHESIS 1: STARTUP ADDS NOISE WHICH INCREASES ROBUSTNESS.

Xie et al. (2020) posit that self-training introduces noise when training the student and thus yielding a more robust student. More robust students may be learning more generalizable representations, and this may be allowing STARTUP to bridge the domain gap. Under this hypothesis, the function of the unlabeled data is only to add noise during training. This in turn suggests that STARTUP should yield improvements on the target tasks *even if trained on unlabeled data from a different domain*. To test this, we train a STARTUP ResNet-18 student on EuroSAT and ImageNet and evaluate it on CropDisease. This model yields a 5-way 1-shot performance of 70.40 ± 0.86 (88.78 ± 0.54 for 5-shot), significantly *underperforming* the naive Transfer baseline (Table 2. See Appendix A.6 for different combinations of unlabeled dataset and target dataset). This suggests that **the hypothesis is incorrect: unlabeled data are not merely functioning as noise**. Rather, STARTUP is learning inherent structure in the target domain useful for downstream classification. The question now becomes what inherent structure STARTUP is learning, which leads us to the next hypothesis.

5.2.2 HYPOTHESIS 2: STARTUP ENHANCES TEACHER-INDUCED GROUPINGS

The teacher produces a meaningful grouping of the data from the target domain. The predictions made by the teacher essentially induce a grouping on the target domain. Even though the base label space and novel label space are disjoint, the groupings produced by the teacher might not be entirely irrelevant for the downstream classification task. To test this, we first assign each example in the novel datasets to its most probable prediction by the teacher (ResNet 18 trained on ImageNet). We then compute the adjusted mutual information (AMI) (Vinh et al., 2010) between the resulting grouping and ground truth label. AMI ranges from 0 for unrelated groupings to 1 for identical groupings. From Table 3, we see that on EuroSAT and CropDisease, there is quite a bit of agreement between the induced grouping and ground truth label. Interestingly, these are the two datasets where we observe the best transfer performance and most improvement from STARTUP (Table 2), suggesting correlations between the agreement and the downstream classification task performance.

STARTUP enhances the grouping induced by the teacher. Even though the induced groupings by the teacher can be meaningful, one could argue that those groupings are captured in the teacher model already, and no further action to update the representation is necessary. However, we posit that STARTUP encourages the feature representations to emphasize the grouping. To verify, we plot the t-SNE (Maaten & Hinton, 2008) of the data prior to STARTUP and after STARTUP for the




Figure 2: t-SNE plot of EuroSAT and CropDisease prior to and after STARTUP.

Table 4: 5-way k-shot classification accuracy on miniImageNet and tieredImageNet.

Methods	miniImageNet		tieredImageNet	
	k=1	k=5	k=1	k=5
Transfer	54.18 ± 0.79	76.20 ± 0.64	57.29 ± 0.83	79.05 ± 0.65
STARTUP-T (no SS)	53.91 ± 0.79	76.26 ± 0.64	60.39 ± 0.86	80.14 ± 0.65
STARTUP (no SS)	53.74 ± 0.80	76.42 ± 0.63	55.49 ± 0.85	78.36 ± 0.66
STARTUP	54.20 ± 0.81	76.48 ± 0.63	55.33 ± 0.85	77.78 ± 0.67

two datasets in figure 2. From the t-SNE plot, we observe more separation after doing STARTUP, signifying a representation with stronger discriminability.

Put together, this suggests that **STARTUP works by (a) inducing a potentially meaningful grouping on the target domain data, and (b) training a representation that emphasizes this grouping.**

5.3 FEW-SHOT TRANSFER ACROSS SIMILAR DOMAINS

Is STARTUP still useful when the gap between the base and target is smaller? To answer this, we tested STARTUP on two popular within-domain few-shot learning benchmark: miniImageNet (Vinyals et al., 2016) and tieredImageNet (Ren et al., 2018). For miniImageNet, we use 20% of the novel set as the unlabeled dataset and use the same teacher as in section 5.1. For tieredImageNet, we use 10% of the novel set as unlabeled data and use ResNet-12 (Oreshkin et al., 2018) as our model architecture. We follow the same evaluation protocols in section 5.1.

We report the results in table 4. We find that on miniImageNet, STARTUP neither helps nor hurts, indicating that the representation is already well-matched. On tieredImageNet, we find that the STARTUP-T variant did outperform naive transfer, but the default STARTUP variant did not. This might be because in this case the pre-trained classifier head does in fact carry useful information for the target domain and therefore should not be discarded. In sum, these results show the potential of STARTUP variants to boost few-shot transfer even when the base and target domains are close.

Additional Ablation Studies: We conducted two additional ablation studies: (a) training the student with various amount of unlabeled data and (b) training the student without the base dataset. We show that STARTUP benefits from more unlabeled data (Appendix A.4) and training student without the base dataset can hurt performance in certain datasets but not all datasets (Appendix A.5).

6 CONCLUSION

We investigate the use of unlabeled data from novel target domains to mitigate the performance degradation of few-shot learners due to large domain/task differences. We introduce STARTUP - a simple yet effective approach that allows few-shot learners to adapt feature representations to the target domain while retaining class grouping induced by the base classifier. We show that STARTUP outperforms prior art on extreme cross-domain few-shot transfer.

7 ACKNOWLEDGEMENT

This work was funded by the DARPA LwLL program.

REFERENCES

- Lars Buitinck, Gilles Louppe, Mathieu Blondel, Fabian Pedregosa, Andreas Mueller, Olivier Grisel, Vlad Niculae, Peter Prettenhofer, Alexandre Gramfort, Jaques Grobler, Robert Layton, Jake VanderPlas, Arnaud Joly, Brian Holt, and Gaël Varoquaux. API design for machine learning software: experiences from the scikit-learn project. In *ECML PKDD Workshop: Languages for Data Mining and Machine Learning*, pp. 108–122, 2013.
- Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. *Proceedings of International Conference on Machine Learning (ICML)*, 2020.
- Wei-Yu Chen, Yen-Cheng Liu, Zsolt Kira, Yu-Chiang Frank Wang, and Jia-Bin Huang. A closer look at few-shot classification. *Proceedings of the International Conference on Learning Representations (ICLR)*, 2019a.
- Zitian Chen, Yanwei Fu, Yu-Xiong Wang, Lin Ma, Wei Liu, and Martial Hebert. Image deformation meta-networks for one-shot learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2019b.
- Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pp. 248–255. Ieee, 2009.
- Guneet S Dhillon, Pratik Chaudhari, Avinash Ravichandran, and Stefano Soatto. A baseline for few-shot image classification. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2020.
- Chelsea Finn, Pieter Abbeel, and Sergey Levine. Model-agnostic meta-learning for fast adaptation of deep networks. In *Proceedings of the 34th International Conference on Machine Learning-Volume 70*, pp. 1126–1135, 2017.
- Chelsea Finn, Kelvin Xu, and Sergey Levine. Probabilistic model-agnostic meta-learning. In *Advances in Neural Information Processing Systems*, pp. 9516–9527, 2018.
- Spyros Gidaris and Nikos Komodakis. Dynamic few-shot visual learning without forgetting. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 4367–4375, 2018.
- Spyros Gidaris and Nikos Komodakis. Generating classification weights with gnn denoising autoencoders for few-shot learning. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 21–30, 2019.
- Spyros Gidaris, Praveer Singh, and Nikos Komodakis. Unsupervised representation learning by predicting image rotations. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2018.
- Spyros Gidaris, Andrei Bursuc, Nikos Komodakis, Patrick Pérez, and Matthieu Cord. Boosting few-shot visual learning with self-supervision. In *Proceedings of the IEEE International Conference on Computer Vision*, pp. 8059–8068, 2019.

- Yunhui Guo, Noel CF Codella, Leonid Karlinsky, John R Smith, Tajana Rosing, and Rogerio Feris. A new benchmark for evaluation of cross-domain few-shot learning. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2020.
- Bharath Hariharan and Ross Girshick. Low-shot visual recognition by shrinking and hallucinating features. In *Proceedings of the IEEE International Conference on Computer Vision*, pp. 3018–3027, 2017.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778, 2016.
- Kaiming He, Haoqi Fan, Yuxin Wu, Saining Xie, and Ross Girshick. Momentum contrast for unsupervised visual representation learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 9729–9738, 2020.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Judy Hoffman, Eric Tzeng, Taesung Park, Jun-Yan Zhu, Phillip Isola, Kate Saenko, Alexei Efros, and Trevor Darrell. Cycada: Cycle-consistent adversarial domain adaptation. In *International conference on machine learning*, pp. 1989–1998. PMLR, 2018.
- Ruibing Hou, Hong Chang, MA Bingpeng, Shiguang Shan, and Xilin Chen. Cross attention network for few-shot classification. In *Advances in Neural Information Processing Systems*, pp. 4003–4014, 2019.
- Issam H Laradji and Reza Babanezhad. M-adda: Unsupervised domain adaptation with deep metric learning. In *Domain Adaptation for Visual Understanding*, pp. 17–31. Springer, 2020.
- Kwonjoon Lee, Subhransu Maji, Avinash Ravichandran, and Stefano Soatto. Meta-learning with differentiable convex optimization. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 10657–10665, 2019.
- Xinzhe Li, Qianru Sun, Yaoyao Liu, Qin Zhou, Shibao Zheng, Tat-Seng Chua, and Bernt Schiele. Learning to self-train for semi-supervised few-shot classification. In *Advances in Neural Information Processing Systems*, pp. 10276–10286, 2019.
- Yanbin Liu, Juho Lee, Minseop Park, Saehoon Kim, Eunho Yang, Sung Ju Hwang, and Yi Yang. Learning to propagate labels: Transductive propagation network for few-shot learning. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2019.
- Mingsheng Long, Zhangjie Cao, Jianmin Wang, and Michael I Jordan. Conditional adversarial domain adaptation. In *Advances in Neural Information Processing Systems*, pp. 1640–1650, 2018.
- Laurens van der Maaten and Geoffrey Hinton. Visualizing data using t-sne. *Journal of machine learning research*, 9(Nov):2579–2605, 2008.
- Ishan Misra and Laurens van der Maaten. Self-supervised learning of pretext-invariant representations. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 6707–6717, 2020.
- Alex Nichol and John Schulman. Reptile: a scalable metalearning algorithm.
- Mehdi Noroozi and Paolo Favaro. Unsupervised learning of visual representations by solving jigsaw puzzles. In *European Conference on Computer Vision*, pp. 69–84. Springer, 2016.
- Boris Oreshkin, Pau Rodríguez López, and Alexandre Lacoste. Tadam: Task dependent adaptive metric for improved few-shot learning. In *Advances in Neural Information Processing Systems*, pp. 721–731, 2018.

- Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Kopf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. Pytorch: An imperative style, high-performance deep learning library. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d’ Alché-Buc, E. Fox, and R. Garnett (eds.), *Advances in Neural Information Processing Systems 32*, pp. 8024–8035. Curran Associates, Inc., 2019. URL <http://papers.neurips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf>.
- Hang Qi, Matthew Brown, and David G Lowe. Low-shot learning with imprinted weights. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 5822–5830, 2018.
- Sachin Ravi and Hugo Larochelle. Optimization as a model for few-shot learning. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2017.
- Mengye Ren, Eleni Triantafillou, Sachin Ravi, Jake Snell, Kevin Swersky, Joshua B Tenenbaum, Hugo Larochelle, and Richard S Zemel. Meta-learning for semi-supervised few-shot classification. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2018.
- Pau Rodríguez, Issam Laradji, Alexandre Drouin, and Alexandre Lacoste. Embedding propagation: Smoother manifold for few-shot classification. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2020.
- Andrei A Rusu, Dushyant Rao, Jakub Sygnowski, Oriol Vinyals, Razvan Pascanu, Simon Osindero, and Raia Hadsell. Meta-learning with latent embedding optimization. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2019.
- Jake Snell, Kevin Swersky, and Richard Zemel. Prototypical networks for few-shot learning. In *Advances in neural information processing systems*, pp. 4077–4087, 2017.
- Jong-Chyi Su, Subhansu Maji, and Bharath Hariharan. When does self-supervision improve few-shot learning? In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2020.
- Qianru Sun, Yaoyao Liu, Tat-Seng Chua, and Bernt Schiele. Meta-transfer learning for few-shot learning. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 403–412, 2019.
- Flood Sung, Yongxin Yang, Li Zhang, Tao Xiang, Philip HS Torr, and Timothy M Hospedales. Learning to compare: Relation network for few-shot learning. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 1199–1208, 2018.
- Hung-Yu Tseng, Hsin-Ying Lee, Jia-Bin Huang, and Ming-Hsuan Yang. Cross-domain few-shot classification via learned feature-wise transformation. In *International Conference on Learning Representations*, 2020.
- Eric Tzeng, Judy Hoffman, Kate Saenko, and Trevor Darrell. Adversarial discriminative domain adaptation. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 7167–7176, 2017.
- Nguyen Xuan Vinh, Julien Epps, and James Bailey. Information theoretic measures for clusterings comparison: Variants, properties, normalization and correction for chance. *The Journal of Machine Learning Research*, 11:2837–2854, 2010.
- Oriol Vinyals, Charles Blundell, Timothy Lillicrap, Daan Wierstra, et al. Matching networks for one shot learning. In *Advances in neural information processing systems*, pp. 3630–3638, 2016.
- Bram Wallace and Bharath Hariharan. Extending and analyzing self-supervised learning across domains. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2020.
- Mei Wang and Weihong Deng. Deep visual domain adaptation: A survey. *Neurocomputing*, 312: 135–153, 2018.

- Yikai Wang, Chengming Xu, Chen Liu, Li Zhang, and Yanwei Fu. Instance credibility inference for few-shot learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 12836–12845, 2020.
- Yu-Xiong Wang, Ross Girshick, Martial Hebert, and Bharath Hariharan. Low-shot learning from imaginary data. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 7278–7286, 2018.
- Garrett Wilson and Diane J Cook. A survey of unsupervised deep domain adaptation. *ACM Transactions on Intelligent Systems and Technology (TIST)*, 11(5):1–46, 2020.
- Zhirong Wu, Yuanjun Xiong, Stella X Yu, and Dahua Lin. Unsupervised feature learning via non-parametric instance discrimination. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 3733–3742, 2018.
- Qizhe Xie, Minh-Thang Luong, Eduard Hovy, and Quoc V Le. Self-training with noisy student improves imagenet classification. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 10687–10698, 2020.
- Ruijia Xu, Guanbin Li, Jihan Yang, and Liang Lin. Larger norm more transferable: An adaptive feature norm approach for unsupervised domain adaptation. In *Proceedings of the IEEE International Conference on Computer Vision*, pp. 1426–1435, 2019.
- I Zeki Yalniz, Hervé Jégou, Kan Chen, Manohar Paluri, and Dhruv Mahajan. Billion-scale semi-supervised learning for image classification. *arXiv preprint arXiv:1905.00546*, 2019.
- Zhongjie Yu, Lin Chen, Zhongwei Cheng, and Jiebo Luo. Transmatch: A transfer-learning scheme for semi-supervised few-shot learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 12856–12864, 2020.
- Richard Zhang, Phillip Isola, and Alexei A Efros. Colorful image colorization. In *European conference on computer vision*, pp. 649–666. Springer, 2016.

A APPENDIX

A.1 IMPLEMENTATION DETAILS

We implemented STARTUP by modifying the the publicly-available implementation¹ of BSCD-FSL by Guo et al. (2020).

A.1.1 TRAINING THE TEACHER

1. MiniImageNet: We train the teacher model using the code provided in the BSCD-FSL benchmark. We keep everything the same except setting the batch size from 16 to 256.
2. TieredImageNet: We used the same setup as miniImageNet except we reduce the number of epochs to 90. We do not use any image augmentation for tieredImageNet.
3. ImageNet: We used the pretrained ResNet18 available on PyTorch (Paszke et al., 2019)

A.1.2 TRAINING THE STUDENT

Optimization Details. Regardless of the base and novel datasets, the student model is trained for 1000 epochs where an epoch is defined to be a complete pass on the unlabeled data. We use a batch size of 256 on the unlabeled dataset and a batch size of 256 for the base dataset if applicable. We use the SGD with momentum optimizer with momentum 0.9 and weight decay $1e-4$. To pick the suitable starting learning rate, 10% of the unlabeled data and 5% of the labeled data (1% when using ImageNet as the base dataset) are set aside as our internal validation set. We pick the starting learning rate by training the student with starting learning rate $lr \in \{1e-1, 5e-2, 3e-2, 1e-2, 5e-3, 3e-3, 1e-3\}$ for k epochs where k is the smallest epoch that guarantees at least 50 updates to the

¹<https://github.com/IBM/cdfsl-benchmark>

model and pick the learning rate that yields lowest loss on the validation set as the starting learning rate. We reduce the learning rate by a factor of 2 when the training loss has not decreased by 20 epochs. The model that achieves the lowest loss on the internal validation set throughout the 1000 epochs of training is picked as the final model.

SimCLR. Our implementation of SimCLR’s loss function is based on a publicly available implementation of SimCLR ². We added the two-layer projection head on top of the embedding function ϕ . The temperature of NT-Xent is set to 1 since there is no validation set for BSCD-FSL for hyperparameter selection and we use a temperature of 1 when inferring the soft label of the unlabeled set. For the stochastic image augmentations for SimCLR, we use the augmentations defined for each novel dataset in Guo et al. (2020). These augmentations include the commonly used “randomly resized crop”, color jittering, random horizontal flipping. For tieredImageNet and miniImageNet, we use the stochastic transformation implemented for the BSCD-FSL benchmark. We refer readers to the BSCD-FSL implementation for more details.

When training the student on the base dataset, we use the augmentation used for training the teacher for fair comparison. The batchsize for SIMCLR is set to 256.

A.1.3 TRAINING LINEAR CLASSIFIER.

We use the implementation by BSCD-FSL, i.e training the linear classifier with standard cross entropy and SGD optimizer. The linear classifier is trained for 100 epochs with learning rate 0.01, momentum 0.9 and weight decay $1e-4$.

A.1.4 BASELINES

We use the same evaluation methods - linear classifier. Please see A.1.3 for classifier training.

Transfer. This is implemented using the teacher model as feature extractor. Please see A.1.1 for details.

SimCLR. This is implemented similarly to the SimCLR loss described in A.1.2

A.1.5 T-SNE

We use the publicly available scikit-learn implementation of t-SNE (Buitinck et al., 2013). We used the default parameters except for the perplexity where we set to 50. To speed up the experiment, we randomly sampled 25% of the data used for sampling few-shot tasks (80 % of the full dataset) and run t-SNE on this subset.

A.2 HIGH SHOTS RESULTS FOR MINIIMAGENET \rightarrow BSCD-FSL

We present the result on miniImageNet \rightarrow BSCD-FSL for higher shots in Table 5 and results on ImageNet \rightarrow BSCD-FSL for higher shots in Table 6. The conclusions we found in 5.1 still hold for higher shots in general.

A.3 RANDOM INITIALIZATION FOR THE STUDENT

We investigate the impact of different initialization strategies for the student on STARTUP. For this experiment, we remove SimCLR from STARTUP and consider three initialization strategies for the student - from scratch (STARTUP-Rand (no SS)), from teacher embedding with a randomly initialized classifier (STARTUP (no SS)), from teacher model (STARTUP-T (no SS)). We repeated the experiment in section 5.1 on miniImageNet \rightarrow BSCD-FSL and report the results in table 7. We found that STARTUP (no SS) and STARTUP-T (no SS) performs similarly whereas STARTUP-Rand (no SS) performs significantly worse than its counterparts on ChestX and ISIC. Interestingly, STARTUP-Rand (no SS) performs better than the rest on CropDisease, suggesting again that the base embedding trained on miniImageNet might not be a good starting point for this dataset as we have found in section 5.1 (comparing SimCLR to STARTUP).

²<https://github.com/sthalles/SimCLR>

Table 5: 5-way k-shot classification accuracy on miniImageNet→BSCD-FSL for higher shots. Mean and 95% confidence interval are reported. * are methods reported in (Guo et al., 2020). Despite using their code, difference in batch size and test set (80% of the original test set) have resulted in discrepancies between our Transfer and their Transfer*. (no SS) indicates removal of SimCLR.

Methods	ChestX		ISIC	
	k=20	k=50	k=20	k=50
MAML*	27.53 \pm 0.43	-	52.36 \pm 0.57	-
ProtoNet*	28.21 \pm 1.15	29.32 \pm 1.12	49.50 \pm 0.55	51.99 \pm 0.52
ProtoNet + FWT*	26.87 \pm 0.43	30.12 \pm 0.46	43.78 \pm 0.47	49.84 \pm 0.51
MetaOpt*	25.53 \pm 1.02	29.35 \pm 0.99	49.42 \pm 0.60	54.80 \pm 0.54
Transfer*	30.83 \pm 1.05	36.04 \pm 0.46	52.78 \pm 0.58	57.34 \pm 0.56
Transfer	31.99 \pm 0.46	35.74 \pm 0.47	54.28 \pm 0.59	60.26 \pm 0.56
SimCLR	29.62 \pm 0.44	32.69 \pm 0.42	47.17 \pm 0.58	52.55 \pm 0.56
STARTUP-T (no SS)	31.77 \pm 0.44	35.57 \pm 0.47	55.80 \pm 0.59	61.15 \pm 0.56
STARTUP (no SS)	33.02 \pm 0.47	36.72 \pm 0.47	57.41 \pm 0.57	62.71 \pm 0.56
STARTUP-T	32.79 \pm 0.46	36.66 \pm 0.47	56.43 \pm 0.60	61.76 \pm 0.58
STARTUP	33.19 \pm 0.46	36.91 \pm 0.50	58.63 \pm 0.58	64.16 \pm 0.58

Methods	EuroSAT		CropDisease	
	k=20	k=50	k=20	k=50
MAML*	81.95 \pm 0.55	-	89.75 \pm 0.42	-
ProtoNet*	82.27 \pm 0.57	80.48 \pm 0.57	88.15 \pm 0.51	90.81 \pm 0.43
ProtoNet + FWT*	75.74 \pm 0.70	78.64 \pm 0.57	85.82 \pm 0.51	87.17 \pm 0.50
MetaOpt*	79.19 \pm 0.62	83.62 \pm 0.58	82.89 \pm 0.54	91.76 \pm 0.38
Transfer*	84.13 \pm 0.52	86.62 \pm 0.47	94.45 \pm 0.36	96.62 \pm 0.25
Transfer	88.31 \pm 0.45	91.09 \pm 0.37	96.10 \pm 0.28	97.69 \pm 0.20
SimCLR	72.25 \pm 0.58	78.64 \pm 0.50	97.26 \pm 0.23	98.44 \pm 0.17
STARTUP-T (no SS)	88.44 \pm 0.46	91.13 \pm 0.38	96.31 \pm 0.28	97.80 \pm 0.20
STARTUP (no SS)	89.29 \pm 0.43	91.94 \pm 0.35	96.91 \pm 0.25	98.20 \pm 0.17
STARTUP-T	88.39 \pm 0.46	91.27 \pm 0.38	96.69 \pm 0.26	98.05 \pm 0.19
STARTUP	89.26 \pm 0.43	91.99 \pm 0.36	97.51 \pm 0.21	98.45 \pm 0.17

A.4 IMPACT OF DIFFERENT AMOUNT OF UNLABELED EXAMPLES

STARTUP uses unlabeled data to adapt feature representations to novel domains. As with all learning techniques, it should perform better with more unlabeled data. To investigate how the amount of unlabeled examples impacts STARTUP, we repeated the miniImageNet \rightarrow ISIC experiments in 5.1 with various amount of unlabeled data (20% of the dataset (2003 examples) is set aside for evaluation). The verdict is clear - STARTUP benefits from more unlabeled data (Figure 3).

A.5 TRAINING THE STUDENT WITHOUT THE BASE DATASET

STARTUP requires joint training on both the base dataset as well as the target domain. But in many cases, the base dataset may not be available. Removing the cross entropy loss on the base dataset when training the student essentially boils down to a fine-tuning paradigm. For miniImageNet \rightarrow BSCD-FSL (Table 8), we find no discernible difference between all datasets except on ISIC where we observe significant degradation in 5-shot performance.

Table 6: 5-way k-shot classification accuracy on ImageNet(ILSVRC 2012)→BSCD-FSL for higher shots.

Methods	ChestX		ISIC	
	k=20	k=50	k=20	k=50
Transfer	30.28 \pm 0.45	32.55 \pm 0.46	55.14 \pm 0.60	60.99 \pm 0.60
STARTUP (no SS)	31.98 \pm 0.47	34.22 \pm 0.47	55.54 \pm 0.57	61.54 \pm 0.55
STARTUP	32.40 \pm 0.45	34.95 \pm 0.48	57.06 \pm 0.58	62.94 \pm 0.56

Methods	EuroSAT		CropDisease	
	k=20	k=50	k=20	k=50
Transfer	91.78 \pm 0.33	93.76 \pm 0.29	96.96 \pm 0.25	98.10 \pm 0.19
STARTUP (no SS)	92.60 \pm 0.31	94.53 \pm 0.26	97.94 \pm 0.20	98.62 \pm 0.16
STARTUP	94.27 \pm 0.26	95.61 \pm 0.23	98.55 \pm 0.17	99.07 \pm 0.13

Table 7: 5-way k-shot classification accuracy on miniImageNet→BSCD-FSL for different initialization strategies. Mean and 95% confidence interval are reported.

Methods	ChestX		ISIC	
	k=1	k=5	k=1	k=5
STARTUP-Rand (no SS)	22.38 \pm 0.41	24.96 \pm 0.41	29.76 \pm 0.60	40.45 \pm 0.59
STARTUP-T (no SS)	22.79 \pm 0.41	26.03 \pm 0.43	32.37 \pm 0.61	45.20 \pm 0.61
STARTUP (no SS)	22.87 \pm 0.41	26.68 \pm 0.45	32.24 \pm 0.62	46.48 \pm 0.61

Methods	EuroSAT		CropDisease	
	k=1	k=5	k=1	k=5
STARTUP-Rand (no SS)	63.44 \pm 0.89	81.05 \pm 0.60	74.44 \pm 0.83	92.04 \pm 0.47
STARTUP-T (no SS)	63.00 \pm 0.84	81.25 \pm 0.62	71.11 \pm 0.83	90.79 \pm 0.49
STARTUP (no SS)	62.90 \pm 0.83	81.81 \pm 0.61	73.30 \pm 0.82	91.69 \pm 0.47

A.6 STARTUP ON DIFFERENT UNLABELED DATA

We consider the ImageNet \rightarrow CD-FSL experiment. We perform STARTUP on unlabeled data different from the target domain and present the result in Table 9. We found that it is crucial that the unlabeled data to perform STARTUP on should be from the target domain of interest.




Figure 3: 5-way 5-shot Classification Accuracy of STARTUP for miniImageNet \rightarrow ISIC with various amount of unlabeled data. Mean and 95% confidence interval over 600 tasks are plotted.

Table 8: 5-way k-shot classification accuracy on miniImageNet \rightarrow BSCD-FSL. We compare STARTUP to fine-tuning.

Methods	ChestX		ISIC	
	k=1	k=5	k=1	k=5
STARTUP	23.09 \pm 0.43	26.94 \pm 0.44	32.66 \pm 0.60	47.22 \pm 0.61
Fine-tuning	22.76 \pm 0.41	27.05 \pm 0.45	32.45 \pm 0.61	45.73 \pm 0.61

Methods	EuroSAT		CropDisease	
	k=1	k=5	k=1	k=5
STARTUP	63.88 \pm 0.84	82.29 \pm 0.60	75.93 \pm 0.80	93.02 \pm 0.45
Fine-tuning	62.86 \pm 0.85	82.36 \pm 0.61	76.13 \pm 0.78	93.01 \pm 0.44

Table 9: Few-shot classification accuracy on ImageNet(ILSVRC 2012) \rightarrow BSCD-FSL with STARTUP on different datasets. STARTUP-X represents the STARTUP student trained on ImageNet and dataset X. The top table presents the results for 5-way 1-shot and the bottom table presents the results for 5-way 5-shot

Methods	ChestX	ISIC	EuroSAT	CropDisease
STARTUP-ChestX	23.03 \pm 0.42	31.02 \pm 0.55	65.20 \pm 0.87	70.36 \pm 0.86
STARTUP-ISIC	21.98 \pm 0.39	31.69 \pm 0.59	61.31 \pm 0.78	69.27 \pm 0.87
STARTUP-EuroSAT	22.23 \pm 0.38	30.38 \pm 0.59	73.83 \pm 0.77	70.40 \pm 0.86
STARTUP-CropDisease	22.51 \pm 0.40	30.59 \pm 0.55	62.56 \pm 0.83	85.10 \pm 0.74

Methods	ChestX	ISIC	EuroSAT	CropDisease
STARTUP-ChestX	27.24 \pm 0.46	44.14 \pm 0.59	83.76 \pm 0.56	89.95 \pm 0.49
STARTUP-ISIC	25.05 \pm 0.42	46.02 \pm 0.59	81.57 \pm 0.55	89.24 \pm 0.50
STARTUP-EuroSAT	25.15 \pm 0.43	43.90 \pm 0.58	89.70 \pm 0.41	88.78 \pm 0.54
STARTUP-CropDisease	25.21 \pm 0.44	44.34 \pm 0.60	83.13 \pm 0.54	96.06 \pm 0.33